

EE3233 Systems Programming for Engrs

Reference: M. Kerrisk, The Linux Programming Interface

Lecture 13

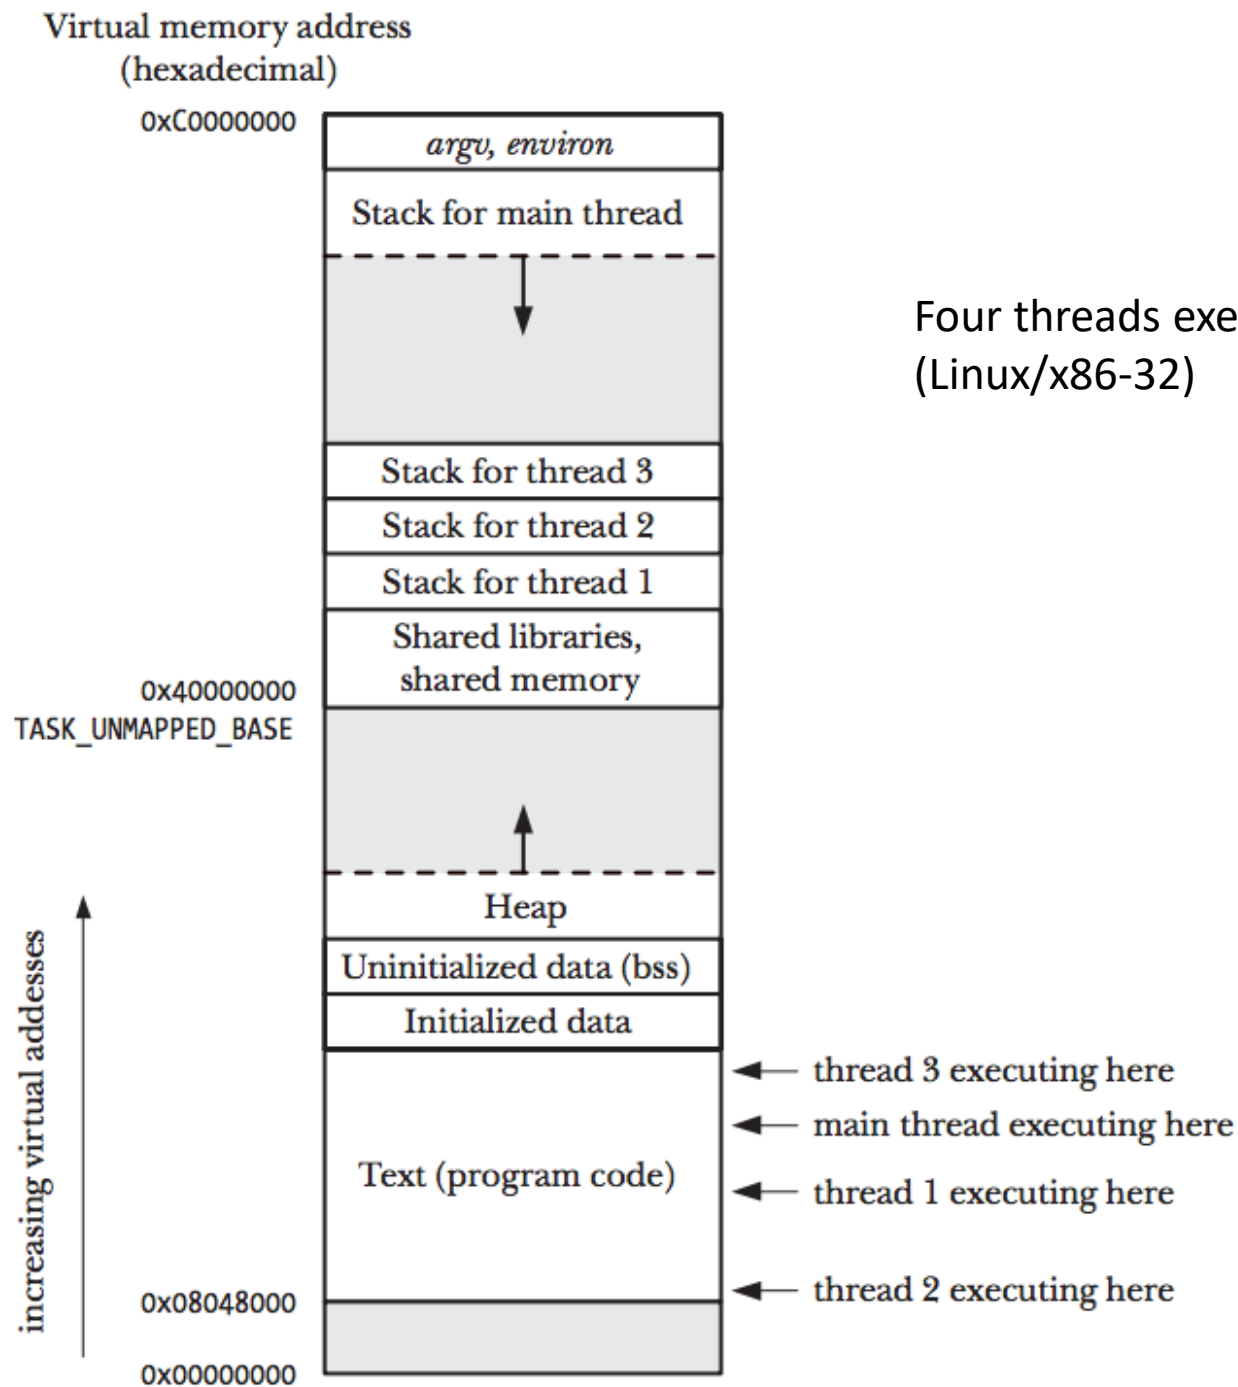
Threads



ECE ELECTRICAL & COMPUTER
ENGINEERING

Thread

- A mechanism that permits an application to perform multiple tasks concurrently
- A single process can contain multiple threads
 - All of these threads are independently executing the same program, and share the same global memory (initialized data, uninitialized data, and heap segments)
 - In a traditional UNIX, a process contains just one thread
 - Threads in a process can execute concurrently
- On a multiprocessor system, multiple threads can execute in parallel
 - If one thread is blocked on I/O, other threads are still eligible to execute



Advantage of Threads over Processes

- In traditional UNIX achieving concurrency by creating multiple processes
 - Difficult to share information between processes: Since the parent and child don't share memory (other than read-only text segment), we must use some form of interprocess communication
 - Process creation with *fork()* is relatively expensive: Even with *copy-on-write* technique, duplicating process attributes such as page table and file descriptor tables is time-consuming

Advantage of Threads over Processes

- Threads address both of problems:
 - Sharing information between threads is easy and fast : In order to avoid updating problem to the same information from multiple threads, synchronization techniques are employed
 - Threads creation is faster than process creation (> x10) : On Linux, threads are implemented using *clone()*. *copy-on-write* or page table duplication is not required

Thread

- Distinct for each thread
 - Thread ID
 - Signal mask
 - Real-time scheduling policy and priority
 - Capabilities (Linux-specific)
 - Stack (local variables and function call linkage info)

POSIX Thread (Pthread)

- Pthreads and types

pthread_t : Thread identifier

pthread_cond_t : Condition variable

pthread_attr_t : Thread attributes object

POSIX Thread (Pthread)

- Return value from Pthread functions
 - return 0 on success and positive value on failure
 - Example program

```
pthread_t *thread;  
int s;  
  
s = pthread_create(&thread, NULL, func, &arg);  
if (s != 0)  
    errExitEN(s, "pthread_create");
```


Thread Creation

```
#include <pthread.h>
```

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start)(void *), void *arg);
```

Returns 0 on success, or a positive error number on error

- calls the function identified by *start* with argument *arg* (i.e., *start(arg)*)
- If we need to pass multiple arguments to *start*, then *arg* can be specified as a pointer to a structure containing the arguments as separate fields
- Thread points to a buffer of type *pthread_t* into which the unique identifier for this thread is copied before *pthread_create* returns

Thread Termination

- Thread's start function performs a return specifying a return value for the thread
- The thread calls *pthread_exit()*
- A thread is canceled using *pthread_cancel()*
- Any of the threads calls *exit()*, or the main thread performs a return, which causes all threads in the process to terminate immediately

pthread_exit()

- terminates calling thread

```
#include <pthread.h>
```

```
void pthread_exit (void *retval);
```

- If the main thread calls *pthread_exit()* instead of calling *exit()* or performing a return, then the other threads continue to execute

Thread IDs

- ID is returned to the caller of *pthread_create()*, and a thread can obtain its own ID using *pthread_self()*

```
#include <pthread.h>
```

```
pthread_t pthread_self (void);
```

- useful for following reasons:
 - Various pthreads functions use thread ID to identify the thread on which they are to act
 - to tag dynamic data structures with the ID of a particular thread

Thread IDs

- to check whether two thread IDs are the same

```
#include <pthread.h>
```

```
int pthread_equal(pthread_t t1, pthread_t t2);
```

- returns nonzero value if *t1* and *t2* are equal, otherwise 0
- Example: check if the ID of the calling thread matches a thread ID saved in the variable *tid*:

```
if ( pthread_equal (tid, pthread_self() )  
    printf("tid matches self\n");
```

Joining with a Terminated Thread

- *pthread_join()* waits for the thread identified by thread to terminate

```
#include <pthread.h>
```

```
int pthread_join (pthread_t thread, void **retval);
```

Joining

```
#include <pthread.h>
#include "tspi_hdr.h"

static void *
threadFunc(void *arg)
{
    char *s = (char *) arg;

    printf("%s", s);

    return (void *) strlen(s);
}
```

```
$ ./simple_thread
Message from main()
Hello world
Thread returned 12
```

```
int
main(int argc, char *argv[])
{
    pthread_t t1;
    void *res;
    int s;

    s = pthread_create(&t1, NULL, threadFunc, "Hello world\n");
    if (s != 0)
        errExitEN(s, "pthread_create");

    printf("Message from main()\n");
    s = pthread_join(t1, &res);
    if (s != 0)
        errExitEN(s, "pthread_join");

    printf("Thread returned %ld\n", (long) res);

    exit(EXIT_SUCCESS);
}
```

threads/simple_thread.c

Thread Synchronization

- Two tools for thread synchronization:
mutexes, condition variables
- mutexes
 - one thread does not try to access a shared variable at the same time as another thread is modifying it
- condition variables
 - inform threads that a shared variable has changed state

Mutexes

- Threads share information via global variables
 - However, multiple threads do not attempt to modify the same variable at the same time
 - One thread should not try to read the value of a variable while another thread is modifying it
- *critical section*
 - is a section of code that accesses a shared resource and whose execution should be *atomic*
 - *atomic*: its execution should not be interrupted by another thread that simultaneously accesses the same shared resource

Following program shows a problem when shared resources are not accessed atomically
threads/thread_incr.c

```
#include <pthread.h>                                     threads/thread_incr.c
#include "tlpi_hdr.h"

static int glob = 0;

static void *                                           /* Loop 'arg' times incrementing 'glob' */
threadFunc(void *arg)
{
    int loops = *((int *) arg);
    int loc, j;

    for (j = 0; j < loops; j++) {
        loc = glob;
        loc++;
        glob = loc;
    }

    return NULL;
}
```

```
int
main(int argc, char *argv[])
{
    pthread_t t1, t2;
    int loops, s;

    loops = (argc > 1) ? getInt(argv[1], GN_GT_0, "num-loops") : 10000000;

    s = pthread_create(&t1, NULL, threadFunc, &loops);
    if (s != 0)
        errExitEN(s, "pthread_create");
    s = pthread_create(&t2, NULL, threadFunc, &loops);
    if (s != 0)
        errExitEN(s, "pthread_create");

    s = pthread_join(t1, NULL);
    if (s != 0)
        errExitEN(s, "pthread_join");
    s = pthread_join(t2, NULL);
    if (s != 0)
        errExitEN(s, "pthread_join");

    printf("glob = %d\n", glob);
    exit(EXIT_SUCCESS);
}
```

Execution Results

```
$ ./thread_incr 1000
```

```
glob = 2000
```

The first thread completed all of its work and terminated before the second thread even started

Execution Results

\$./thread_incr 10000000

glob = 12039302

Why not 200000000?

Current
value of
glob

2000

3000

2001

Thread 1

Repeatedly:
loc = glob;
loc++;
glob = loc;

loc = glob;

time slice
expires

time slice
begins

Thread 2

Repeatedly:
loc = glob;
loc++;
glob = loc;

time slice
begins

time slice
ends

loc++;
glob = loc;

Key

Executing
on CPU


Waiting
for CPU

Execution Results

- Thread 1 (*t1*) fetches the current value of *glob* into *loc* (assume *glob* = 2000)
- Scheduler time slice for *t1* is expired, and *t2* commences execution
- *t2* performs multiple loops in which it fetches the current value of *glob* into *loc* (starts from *glob*=2000, when time slice terminates for *t2*, assume *glob*=3000)
- *t1* receives another time slice and resumes. It now increases *loc* and assigns to *glob* (2001)
→ effect of increase by *t2* is lost (3000 → 2001)

Eliminating the problem

- Three statements inside the for loop → a single statement:



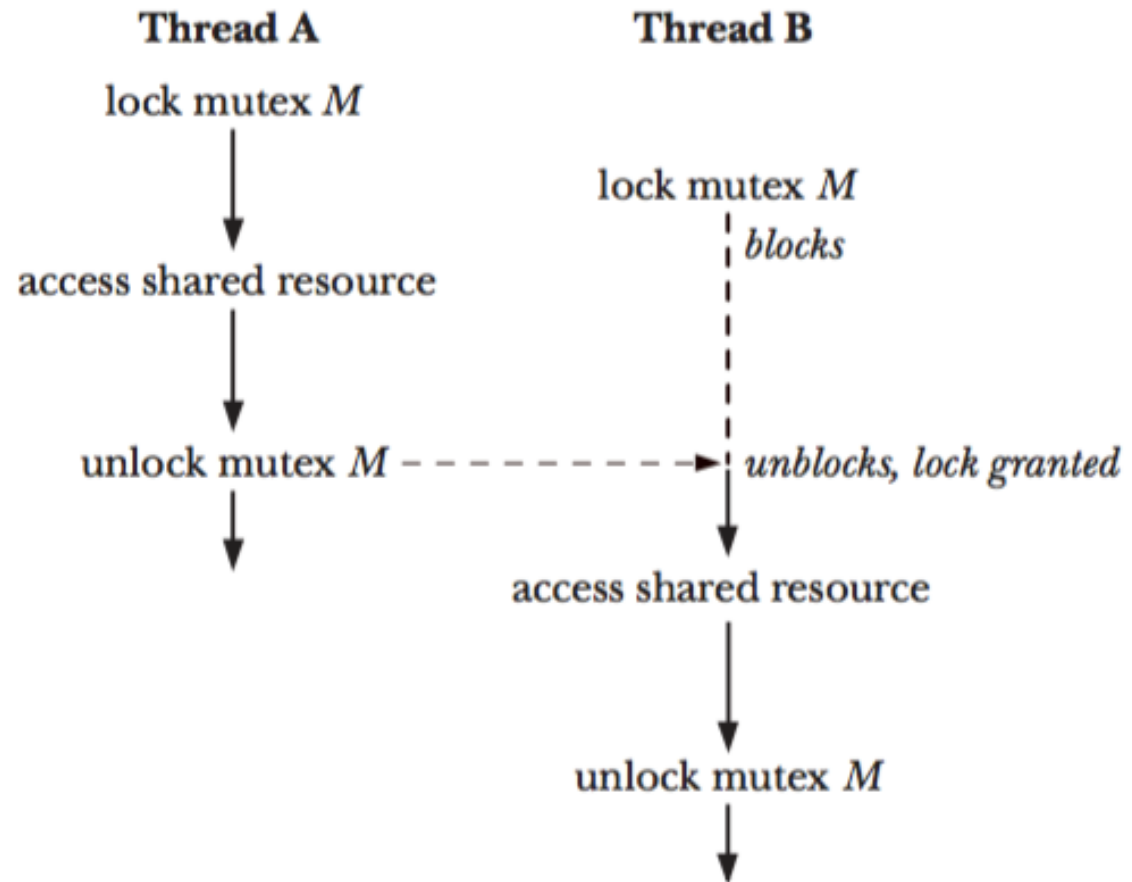
- However, on many H/W architecture, the compiler would still need to convert this single statement into machine code whose steps are equivalent to the three statements
 - To avoid the problems, use a *mutex* (mutual exclusion)

Mutex

- has two states
 - locked
 - unlocked
- At any moment, at most one thread may hold the lock on a mutex
 - Attempting to lock a mutex that is already locked either blocks or fails with an error
- When a thread locks a mutex, it becomes owner of that mutex
 - Only the owner can unlock

Mutex

- Protocol for accessing a resource
 - lock the mutex for the shared resource
 - access the shared resource
 - unlock the mutex



Statically Allocated Mutexes

- A mutex can either be
 - allocated as a static variable
 - created dynamically at run time
- variable type of mutex: `pthread_mutex_t`
- **Before using, a mutex must always be initialized**

```
pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
```

Locking and Unlocking a Mutex

```
#include <pthread.h>
```

```
int pthread_mutex_lock (pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock (pthread_mutex_t *mutex);
```

both return 0 on success, or a positive error number on error

- To lock a mutex, call *pthread_mutex_lock()*
 - If the mutex is currently unlocked, this call locks the mutex and returns immediately
 - If the mutex is currently locked by another thread, then this blocks until the mutex is unlocked

Locking and Unlocking a Mutex

```
#include <pthread.h>
```

```
int pthread_mutex_lock (pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock (pthread_mutex_t *mutex);
```

both return 0 on success, or a positive error number on error

- *pthread_mutex_unlock()* unlocks a mutex previously locked by the calling thread
- It is error
 - to unlock a mutex that is not currently locked
 - to unlock a mutex that is locked by another thread

```
#include <pthread.h>
#include "tlpi_hdr.h"
```

```
static int glob = 0;
static pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;
```

```
static void * /* Loop 'arg' times incrementing 'glob' */
```

```
threadFunc(void *arg)
```

```
{
```

```
    int loops = *((int *) arg);
```

```
    int loc, j, s;
```

```
    for (j = 0; j < loops; j++) {
```

```
        s = pthread_mutex_lock(&mtx);
```

```
        if (s != 0)
```

```
            errExitEN(s, "pthread mutex lock");
```

```
        loc = glob;
```

```
        loc++;
```

```
        glob = loc;
```

```
        s = pthread_mutex_unlock(&mtx);
```

```
        if (s != 0)
```

```
            errExitEN(s, "pthread_mutex_unlock");
```

```
    }
```

```
    return NULL;
```

```
}
```

```
int
main(int argc, char *argv[])
{
    pthread_t t1, t2;
    int loops, s;

    loops = (argc > 1) ? getInt(argv[1], GN_GT_0, "num-loops") : 10000000;

    s = pthread_create(&t1, NULL, threadFunc, &loops);
    if (s != 0)
        errExitEN(s, "pthread_create");
    s = pthread_create(&t2, NULL, threadFunc, &loops);
    if (s != 0)
        errExitEN(s, "pthread_create");

    s = pthread_join(t1, NULL);
    if (s != 0)
        errExitEN(s, "pthread_join");
    s = pthread_join(t2, NULL);
    if (s != 0)
        errExitEN(s, "pthread_join");

    printf("glob = %d\n", glob);
    exit(EXIT_SUCCESS);
}
```

Execution Results

```
$ ./thread_incr_mutex 10000000  
glob = 20000000
```


Mutex Deadlocks

- Sometimes, a thread needs to simultaneously access two or more different shared resources
 - each of which is governed by a separate mutex
- When more than one thread is locking the same set of mutexes, deadlock situations can arise

Thread A	Thread B
1. <i>pthread_mutex_lock(mutex1);</i>	1. <i>pthread_mutex_lock(mutex2);</i>
2. <i>pthread_mutex_lock(mutex2);</i> blocks	2. <i>pthread_mutex_lock(mutex1);</i> blocks

- To avoid such deadlocks, define a mutex hierarchy
 - When threads can lock the same set of mutexes, they should always lock them in the same order
 - Two threads always lock the mutexes in the order of mutex1 followed by mutex2

Dynamic Initialization

```
#include <pthread.h>
```

```
int pthread_mutex_init(pthread_mutex_t *mutex, const pthread_mutexattr_t *attr);
```

```
return 0 on success, or a positive error number on error
```

- *mutex* : to be initialized
- *attr* : is a pointer to a *pthread_mutexattr_t* to define attributes for the mutex (If it is NULL, the mutex is assigned default attributes)
- is required in the following cases:
 - the mutex was dynamically allocated on the heap (linked list contains a *pthread_mutex_t* field)
 - the mutex is an automatic variable allocated on the stack
 - we want to initialize a statically allocated mutex with attributes other than the defaults

Condition Variables

- allows one thread to inform other threads about changes in the state of a shared variable and allows the other threads to wait (block) for such notification

Without Condition Variables

- check “threads/prod_no_condvar.c”
- Producer(main) wastes CPU time due to continuous loop checking the state of the variable *avail*
- condition variable remedies this problem
 - it allows a thread to sleep until another thread notifies (signals) it

Static Allocation of Condition Variables

```
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
```

- As with a mutex, a condition variable must be initialized before use

Signaling and Waiting on Condition Variables

```
#include <pthread.h>
```

```
int pthread_cond_signal (pthread_cond_t *cond);
```

```
int pthread_cond_wait (pthread_cond_t *cond, pthread_mutex_t *mutex);
```

- condition variable always has an associated mutex
 - Both of these objects are passed as arguments to *pthread_cond_wait()*
- *pthread_cond_wait()* performs the following steps:
 - unlock the mutex specified by *mutex* (so that other threads can access the shared variable)
 - block the calling thread until another thread signals the condition variable *cond*; and
 - relock *mutex* (since the thread then immediately accesses the shared variable)

Signaling and Waiting on Condition Variables

- check “threads/prod_condvar.c”

Signaling and Waiting on Condition Variables

```
static pthread_mutex_t mtx = PTHREAD_MUTEX_INITIALIZER;  
static pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
```

```
static void *  
threadFunc(void *arg) {
```

```
    int cnt = atoi((char *) arg);  
    int s, j;
```

```
    for (j=0; j < cnt; j++) {  
        sleep (1);  
        s = pthread_mutex_lock(&mtx);  
        if (s != 0)  
            errExitEN(s, "pthread_mutex_lock");  
        avail++;  
        s = pthread_mutex_unlock(&mtx);  
        if (s != 0)  
            errExitEN(s, "pthread_mutex_unlock");  
        s = pthread_cond_signal(&cond);  
        if (s != 0)  
            errExitEN(s, "pthread_cond_signal");  
    }
```

```
    return NULL;  
}
```


Signaling and Waiting on Condition Variables

```
for (;;) {  
    s = pthread_mutex_lock(&mtx);                                     /* main thread : consumer */  
    if (s != 0)  
        errExitEN(s, "pthread_mutex_lock");  
  
    while (avail == 0) {                                             /* Wait for something to consume */  
        s = pthread_cond_wait(&cond, &mtx);  
        if (s != 0)  
            errExitEN(s, "pthread_cond_wait");  
    }  
  
    while (avail > 0) {                                             /* Consume all available units */  
        /* Do something with produced unit */  
        avail--;  
    }  
  
    s = pthread_mutex_unlock(&mtx);  
    if (s != 0)  
        errExitEN(s, "pthread_mutex_unlock");  
  
    /* Perhaps do other work here that doesn't require mutex lock */  
}
```