

**EE3233 Systems Programming for Engrs**

Reference: M. Kerrisk, The Linux Programming Interface

# Lecture 12

## Program Execution



**ECE** ELECTRICAL & COMPUTER  
**ENGINEERING**

# Executing a New Program: `execve()`

- *execve()* system call loads a new program into a process's memory
  - The old program is discarded, and the process's stack, data, and heap are replaced by new program
  - New program commences execution at its *main()* function
- The most frequent use of *execve()* is in the child produced by a *fork()*
- Various library functions (beginning with `exec`) are layered on top of the *execve()* system call

# Executing a New Program: `execve()`

```
#include <unistd.h>
```

```
int execve (const char *pathname, char *const argv[], char *const envp[]);  
           Never returns on success; return -1 on error
```

- *pathname*: pathname of the new program
- *argv*: command line arguments to be passed to the new program
- *envp*: environment list (name=value)
- After `execve()`, the process ID remains the same

# Executing a New Program: `execve()`

```
#include "tldpi_hdr.h"
```

`procexec/t_execve.c`

```
int
main(int argc, char *argv[])
{
    char *argVec[10];          /* Larger than required */
    char *envVec[] = { "GREET=salut", "BYE=adieu", NULL };

    if (argc != 2 || strcmp(argv[1], "--help") == 0)
        usageErr("%s pathname\n", argv[0]);

    argVec[0] = strrchr(argv[1], '/');      /* Get basename from argv[1] */
    if (argVec[0] != NULL)
        argVec[0]++;
    else
        argVec[0] = argv[1];
    argVec[1] = "hello world";
    argVec[2] = "goodbye";
    argVec[3] = NULL;              /* List must be NULL-terminated */

    execve(argv[1], argVec, envVec);
    errExit("execve");           /* If we get here, something went wrong */
}
```

# Executing a New Program: `execve()`

```
#include "tlpi_hdr.h"
```

`procexec/envargs.c`

```
extern char **environ;
```

```
int
```

```
main(int argc, char *argv[])
```

```
{
```

```
    int j;
```

```
    char **ep;
```

```
    for (j = 0; j < argc; j++)
```

```
        printf("argv[%d] = %s\n", j, argv[j]);
```

```
    for (ep = environ; *ep != NULL; ep++)
```

```
        printf("environ: %s\n", *ep);
```

```
    exit(EXIT_SUCCESS);
```

```
}
```

```
$ ./t_execve ./envargs
argv[0] = envargs
argv[1] = hello world
argv[2] = goodbye
environ: GREET=salut
environ: BYE=adieu
```

# exec()

Function	Specification of program file ( $\neg$ , $p$ )	Specification of arguments ( $v$ , $l$ )	Source of environment ( $e$ , $\neg$ )
<i>execve()</i>	pathname	array	<i>envp</i> argument
<i>execle()</i>	pathname	list	<i>envp</i> argument
<i>execlp()</i>	filename + PATH	list	caller's <i>environ</i>
<i>execvp()</i>	filename + PATH	array	caller's <i>environ</i>
<i>execv()</i>	pathname	array	caller's <i>environ</i>
<i>execl()</i>	pathname	list	caller's <i>environ</i>

# execlp()

procexec/t\_execlp.c

```
#include "t1pi_hdr.h"

int
main(int argc, char *argv[])
{
    if (argc != 2 || strcmp(argv[1], "--help") == 0)
        usageErr("%s pathname\n", argv[0]);

    execlp(argv[1], argv[1], "hello world", (char *) NULL);
    errExit("execlp");          /* If we get here, something went wrong */
}
```

```
$ PATH=/home/mtk/bin:/usr/local/bin:/usr/bin
$ ./t_execlp echo
ERROR [ENOENT No such file or directory] execlp
$ ./t_execlp /bin/echo
hello world
```

# Interpreter Scripts

- An interpreter is a program that reads commands in text form and executes them
  - *awk, sed, perl, python, ruby*
- UNIX kernels allow interpreter scripts to be run in the same way as a binary program file if
  - execute permission is enabled (`chmod +x <file>`)
  - file contains an initial line that specifies the pathname of the interpreter to be used to run the script

***#!/ interpreter-path [optional-arg]***



# Execution of Interpreter Scripts

- When *execve()* is used to run the script, *execve()* detects that the file has 2-byte sequence **#!**, then it extracts the remainder of the line (*pathname + argument*), and execs the interpreter file with following list of arguments

***interpreter-path [optional-arg] script-path arg ...***

# Argument list given to interpreter

**Script file** (located at *script-path*)

*#! interpreter-path optional-arg*

**execve() call within program**  
*execve(script-path, argv, envp)*

*interpreter-path optional-arg script-path arg...*

**Argument list given to interpreter**

# Argument list given to interpreter

```
$ cat > necho.script  
#!/home/mtk/bin/necho some argument  
Some junk
```

*Type Control-D*

```
$ chmod +x necho.script
```

```
$ ./t_execve necho.script
```

```
argv[0] = /home/mtk/bin/necho  
argv[1] = some argument  
argv[2] = necho.script  
argv[3] = hello world  
argv[4] = goodbye
```

*Create script*

*Make script executable*

*And exec the script*

*First 3 arguments are generated by kernel  
Script argument is treated as a single word  
This is the script path  
This was argv[1] given to execve()  
And this was argv[2]*

# /proc/necho.c

```
#include "tlpi_hdr.h"

int
main(int argc, char *argv[])
{
    int j;

    for (j = 0; j < argc; j++)
        printf("argv[%d] = %s\n", j, argv[j]);

    exit(EXIT_SUCCESS);
}
```

# Be careful!!!

```
#include <string.h>
#include <stdlib.h>
#include <stdio.h>
enum { BUFFERSIZE = 512};
void func(const char *input){
    char cmdbuf[BUFFERSIZE];
    int len_wanted = snprintf(cmdbuf, BUFFERSIZE,
                             "any_cmd '%s'", input);
    if (len_wanted >= BUFFERSIZE) {
        /* Handle error */
    } else if (len_wanted < 0) {
        /* Handle error */
    } else if (system(cmdbuf) == -1) {
        /* Handle error */
    }
}
```

<https://wiki.sei.cmu.edu/confluence/pages/viewpage.action?pageId=87152177>