

EE3233 Systems Programming for Engrs

Reference: M. Kerrisk, The Linux Programming Interface

Lecture 16

System V Shared Memory



ECE ELECTRICAL & COMPUTER
ENGINEERING

Shared Memory

- allows two or more processes to share the same region of physical memory
- Shared memory becomes part of a process's user-space memory
 - no kernel intervention is required
 - Some method of synchronization is required so that processes don't simultaneously access the shared memory
 - We have learned System V semaphore as a natural method for such synchronization

Overview

- Typical steps of using a shared memory
 - ① **shmget()**: create a new shared memory or obtain identifier of an existing segment
 - ② **shmat()**: attach the shared memory (make the part of the virtual memory of the calling process)
 - ③ use shared memory referring *addr* returned by *shmat()*
 - ④ **shmdt()**: detach the shared memory segment
 - ⑤ **shmctl()**: delete the shared memory segment

Creating or Opening a Shared Memory

```
#include <sys/types.h>
```

```
#include <sys/shm.h>
```

```
int shmget(key_t key, size_t size, int shmflg);
```

returns shared memory segment identifier on success, or -1 on error

- creates a new shared memory segment or obtains the identifier of an existing segment
- The content of a newly created shared memory segment are initialized to 0
- *key*: unique key. **IPC_PRIVATE** always results in the creation of a new object guaranteed to have a unique identifier
- *size*: desired size of the segment in bytes
- *shmflg*: permissions to be placed or checked. **IPC_CREAT** creates a new segment

Using Shared Memory

```
#include <sys/types.h>
```

```
#include <sys/shm.h>
```

```
int *shmat(int shmid, const void *shmaddr, int shmflg);
```

returns address at which shared memory is attached on success, or -1 on error

- attaches the shared memory segment identified by *shmid* to the calling process's virtual address space
- *shmaddr*
 - if NULL, then segment is attached at a suitable address selected by the kernel
 - if not NULL, and **SHM_RND** is not set, then the segment is attached at the address specified by *shmaddr*
 - if not NULL, and **SHM_RND** is set, then the segment is mapped at the address provided in *shmaddr* rounded down to the nearest multiple of the constant **SHMLBA**(SHared Memory Low Boundary Address)
- *shmflg*
 - To attach a shared memory segment for read-only access, specify **SHM_RDONLY** in *shmflg*
 - If **SHM_RDONLY** is not specified, memory can be both read and modified

Using Shared Memory

```
#include <sys/types.h>
```

```
#include <sys/shm.h>
```

```
int *shmdt(const void *shmaddr);
```

returns 0 on success, or -1 on error

- When a process no longer needs to access a shared memory segment, call *shmdt()* to detach
- *shmaddr*
 - identifies the segment to be detached

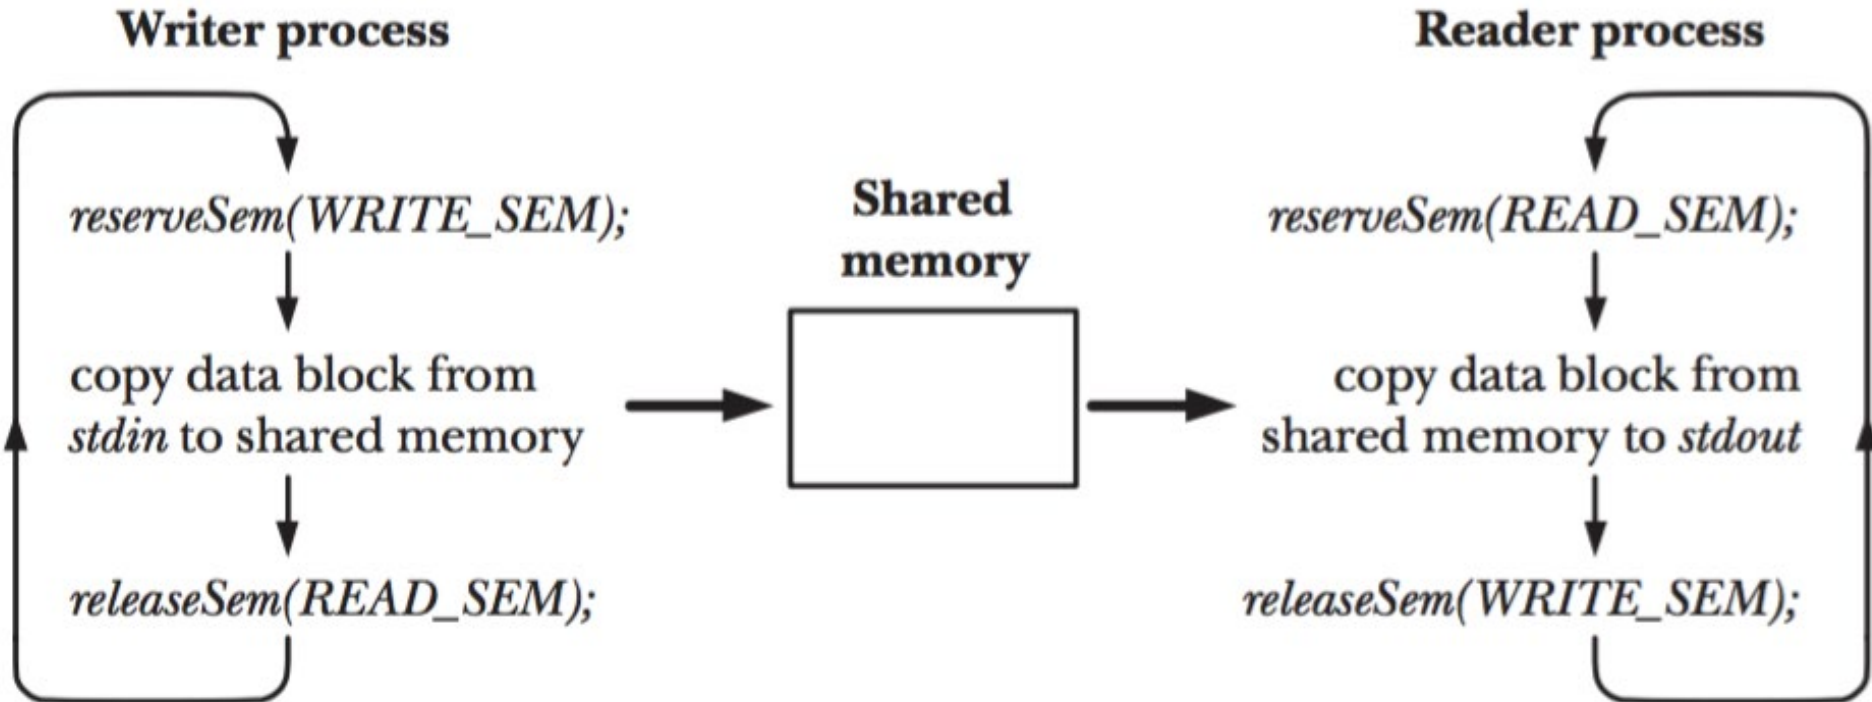
deleting

- use *shmctl()* **IPC_RMID** operation
- During an *exec()*, all attached shared memory segments are detached,
- Shared memory segments are also automatically detached on process termination

Transferring Data via Shared Memory

- [Example] We use shared memory and semaphores
 - consists of two programs: *writer* and *reader*
- Two programs employ a pair of semaphores in a binary semaphore protocol
 - *initSemAvailable()*
 - *initSemInUse()*
 - *reserveSem()*
 - *releaseSem()*

Transferring Data via Shared Memory



```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/sem.h>
#include <sys/shm.h>
#include "binary_sems.h"      /* Declares our binary semaphore functions */
#include "tlpi_hdr.h"

#define SHM_KEY 0x1234        /* Key for shared memory segment */
#define SEM_KEY 0x5678        /* Key for semaphore set */

#define OBJ_PERMS (S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP)
                        /* Permissions for our IPC objects */

#define WRITE_SEM 0           /* Writer has access to shared memory */
#define READ_SEM 1           /* Reader has access to shared memory */

#ifndef BUF_SIZE               /* Allow "cc -D" to override definition */
#define BUF_SIZE 1024         /* Size of transfer buffer */
#endif

struct shmseg {               /* Defines structure of shared memory segment */
    int cnt;                  /* Number of bytes used in 'buf' */
    char buf[BUF_SIZE];       /* Data being transferred */
};
```

```
#include "semun.h"
#include "svshm_xfr.h"
```

```
/* Definition of semun union */
```

[svshm/svshm_xfr_writer.c](#)

```
int
main(int argc, char *argv[])
{
    int semid, shmid, bytes, xfrs;
    struct shmseg *shmp;
    union semun dummy;
```

```
①    semid = semget(SEM_KEY, 2, IPC_CREAT | OBJ_PERMS);
    if (semid == -1)
        errExit("semget");
    if (initSemAvailable(semid, WRITE_SEM) == -1)
        errExit("initSemAvailable");
    if (initSemInUse(semid, READ_SEM) == -1)
        errExit("initSemInUse");
```


```
②    shmid = shmget(SHM_KEY, sizeof(struct shmseg), IPC_CREAT | OBJ_PERMS);
    if (shmid == -1)
        errExit("shmget");

    shmp = shmat(shmid, NULL, 0);
    if (shmp == (void *) -1)
        errExit("shmat");
```

```
/* Transfer blocks of data from stdin to shared memory */
```

```
③ for (xfrs = 0, bytes = 0; ; xfrs++, bytes += shmp->cnt) {
```

```
④   if (reserveSem(semid, WRITE_SEM) == -1)           /* Wait for our turn */  
       errExit("reserveSem");
```



```
⑤   shmp->cnt = read(STDIN_FILENO, shmp->buf, BUF_SIZE);  
   if (shmp->cnt == -1)  
       errExit("read");
```

```
⑥   if (releaseSem(semid, READ_SEM) == -1)           /* Give reader a turn */  
       errExit("releaseSem");
```

```
/* Have we reached EOF? We test this after giving the reader  
   a turn so that it can see the 0 value in shmp->cnt. */
```

```
⑦   if (shmp->cnt == 0)  
       break;  
}
```

```
/* Wait until reader has let us have one more turn. We then know  
   reader has finished, and so we can delete the IPC objects. */
```

```
⑧ if (reserveSem(semid, WRITE_SEM) == -1)  
    errExit("reserveSem");  
  
⑨ if (semctl(semid, 0, IPC_RMID, dummy) == -1)  
    errExit("semctl");  
if (shmdt(shmp) == -1)  
    errExit("shmdt");  
if (shmctl(shmid, IPC_RMID, 0) == -1)  
    errExit("shmctl");  
  
fprintf(stderr, "Sent %d bytes (%d xfrs)\n", bytes, xfrs);  
exit(EXIT_SUCCESS);  
}
```

```
#include "svshm_xfr.h"
```

svshm/svshm_xfr_reader.c

```
int  
main(int argc, char *argv[])  
{  
    int semid, shmid, xfrs, bytes;  
    struct shmseg *shmp;  
  
    /* Get IDs for semaphore set and shared memory created by writer */
```

```
①    semid = semget(SEM_KEY, 0, 0);  
    if (semid == -1)  
        errExit("semget");
```

```
    shmid = shmget(SHM_KEY, 0, 0);  
    if (shmid == -1)  
        errExit("shmget");
```

```
②    shmp = shmat(shmid, NULL, SHM_RDONLY);  
    if (shmp == (void *) -1)  
        errExit("shmat");
```



```
/* Transfer blocks of data from shared memory to stdout */
```

[svshm/svshm_xfr_reader.c](#)

```
③ for (xfrs = 0, bytes = 0; ; xfrs++) {
④     if (reserveSem(semid, READ_SEM) == -1)          /* Wait for our turn */
        errExit("reserveSem");

⑤     if (shmp->cnt == 0)                             /* Writer encountered EOF */
        break;
        bytes += shmp->cnt;

⑥     if (write(STDOUT_FILENO, shmp->buf, shmp->cnt) != shmp->cnt)
        fatal("partial/failed write");


⑦     if (releaseSem(semid, WRITE_SEM) == -1)         /* Give writer a turn */
        errExit("releaseSem");
}

⑧ if (shmdt(shmp) == -1)
    errExit("shmdt");

/* Give writer one more turn, so it can clean up */

⑨ if (releaseSem(semid, WRITE_SEM) == -1)
    errExit("releaseSem");

fprintf(stderr, "Received %d bytes (%d xfrs)\n", bytes, xfrs);
exit(EXIT_SUCCESS);
}
```



Results

```
$ wc -c /etc/services
```

```
764360 /etc/services
```

```
$ ./svshm_xfr_writer < /etc/services &
```

```
[1] 9403
```

```
$ ./svshm_xfr_reader > out.txt
```

```
Received 764360 bytes (747 xfrs)
```

```
Sent 764360 bytes (747 xfrs)
```

```
[1]+  Done                  ./svshm_xfr_writer < /etc/services
```

```
$ diff /etc/services out.txt
```

```
$
```

Display size of test file

Message from reader

Message from writer

Transferring Data via Shared Memory

