

## **EE3233 Systems Programming for Engrs**

Reference: M. Kerrisk, The Linux Programming Interface

# **Lecture 18**

# **Sockets**



**ECE** ELECTRICAL & COMPUTER  
**ENGINEERING**

# Overview

- A method of IPC that exchanges data between applications either on the same host or on different host connected via a network
- Applications communicate using sockets as follows:
  - Each application creates a socket (apparatus) that allows communication, and both applications require one
  - The server binds its socket to a well-known address so that clients can locate it

# Socket System Calls

- ***socket()***
  - creates a new socket
- ***bind()***
  - binds a socket to an address
- ***listen()***
  - allows a stream to accept incoming connections from other sockets
- ***accept()***
  - accepts a connection from a peer application
- ***connect()***
  - establishes a connection with another socket

# socket()

```
fd = socket (domain, type, protocol);
```

- Communication domains, which determines:
  - method of identifying a socket (format of a socket “address”)
  - range of communication (on the same host or different hosts)
- Modern OS supports following domains:
  - **AF\_UNIX**: allows communication between applications on the same host (address structure: *sockaddr\_un*)
  - **AF\_INET**: allows communication between applications running on hosts connected via an Internet Protocol version 4 (IPv4) network (address structure: *sockaddr\_in*)
  - **AF\_INET6**: for IPv6 (address structure: *sockaddr\_in6*)

# socket()

```
fd = socket (domain, type, protocol);
```

- Socket type

- stream (**SOCK\_STREAM**)

- ① Reliable: the transmitted data will arrive at the receiving application exactly as it was transmitted by sender
    - ② Bidirectional: transmitted in either direction
    - ③ Byte-stream: there is no concept of message boundaries
    - ④ Connection-Oriented: operates in connected pair
  - : TCP

- datagram (**SOCK\_DGRAM**)

- ① Data transmission is not reliable
    - ② Message may arrive out-of-order
    - ③ Message boundaries are preserved
    - ④ Connection-Less
  - : UDP

# socket()

```
fd = socket (domain, type, protocol);
```

- protocol values are always specified as 0 for the socket types that we describe

# bind()

```
bind (int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- *sockfd*
  - file descriptor obtained from a previous call to *socket()*
- *addr*
  - a pointer to a structure specifying the address to which this socket is to be bound
- *addrlen*
  - size of the address structure

# bind()

```
int bind (int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

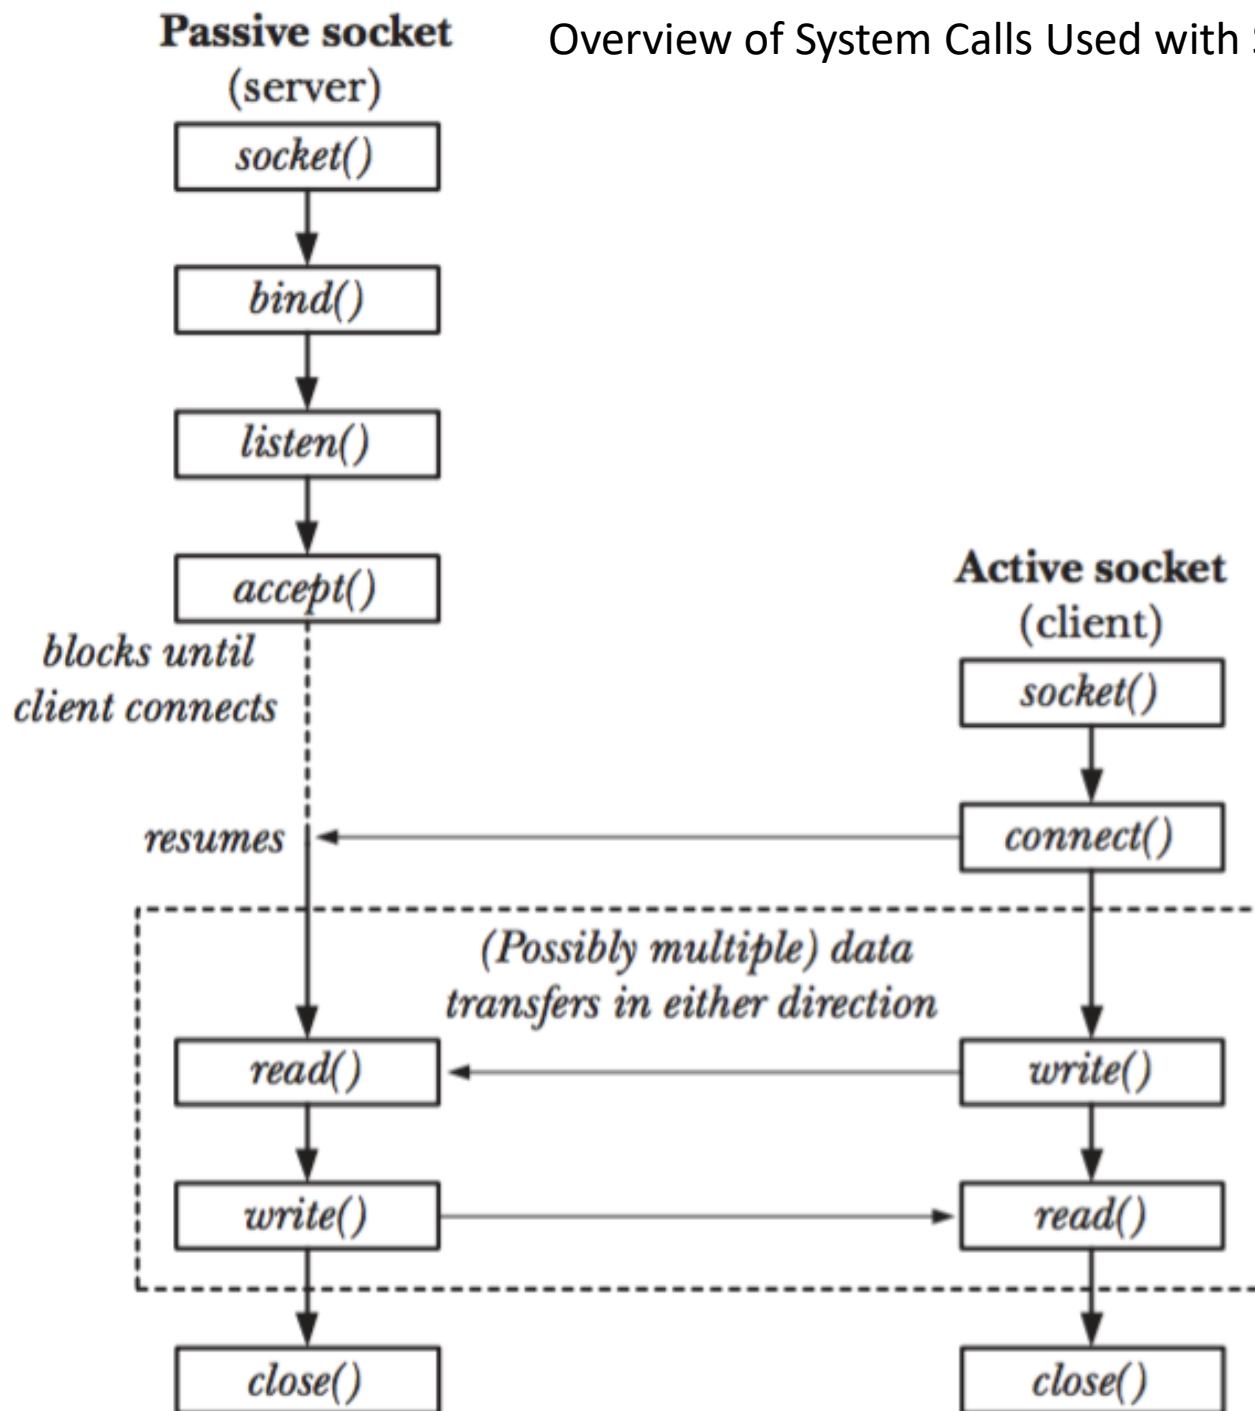
- Each socket domain uses a different address format
  - UNIX domain sockets use pathnames
  - Internet domain sockets use IP + port
- socket API defines a generic address structure

```
struct sockaddr {  
    sa_family_t sa_family;    /* Address family */  
    char        sa_data[14];  /* Socket address */  
};
```



# Stream Sockets

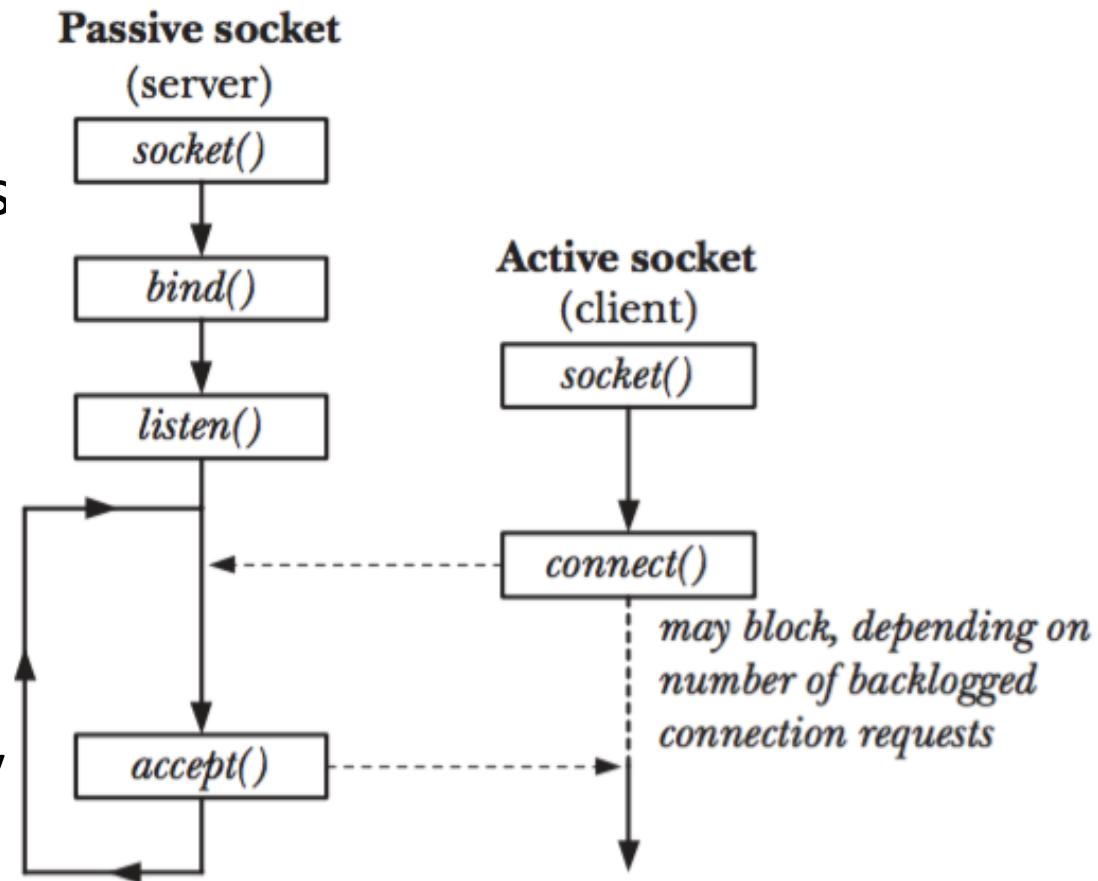
- analogous to a telephone call:
  - ***socket()*** is installing a telephone(socket). For two persons(applications), each of telephone must be installed
  - ***bind()*** is having a known telephone number
  - ***listen()*** is ensuring that our telephone is turned on so that people can call us
  - ***accept()*** is picking up the phone when it rings
  - ***connect()*** is dialing someone's telephone number



# listen()

```
int listen (int sockfd, int backlog);
```

- notifying kernel of its willingness to accept incoming connections
- *sockfd*: a file descriptor
- *backlog*: number of pending connections
  - Connection requests up to this limit succeed immediately



# accept()

```
int accept (int sockfd, struct sockaddr *addr, socklen_t addrlen);
```

returns *file descriptor* on success, or -1 on error

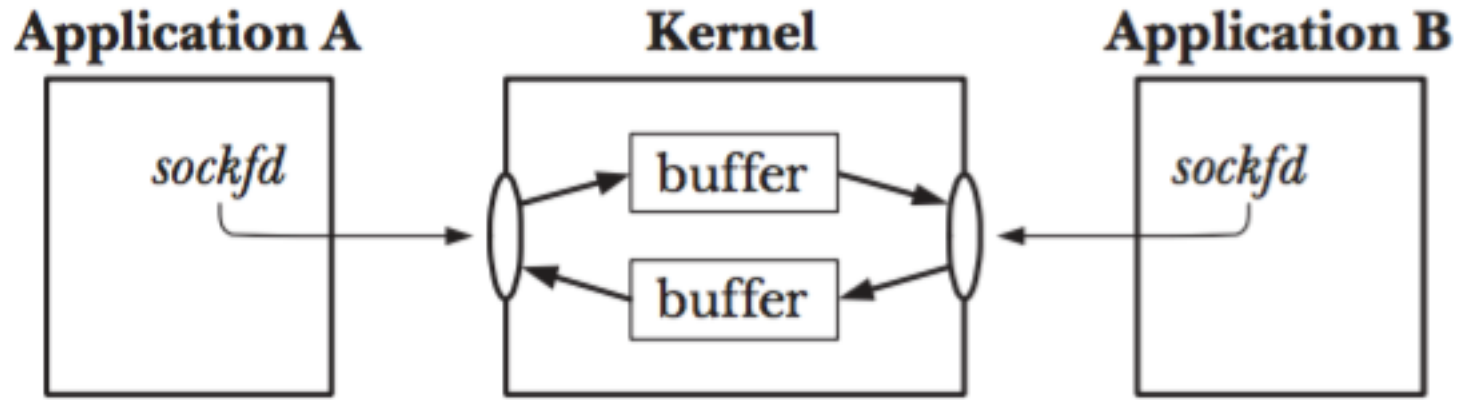
- *sockfd*: a file descriptor
- *addr*: returned address of the peer socket
- *addrlen*: that must be initialized (0) (prior to call) to the size of the buffer pointed to by *addr*

# connect()

```
int connect (int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

- If *connect()* fails and we wish to reattempt the connection, then close the socket, create a new socket, and reattempt the connection with the new socket

# I/O on Stream Sockets

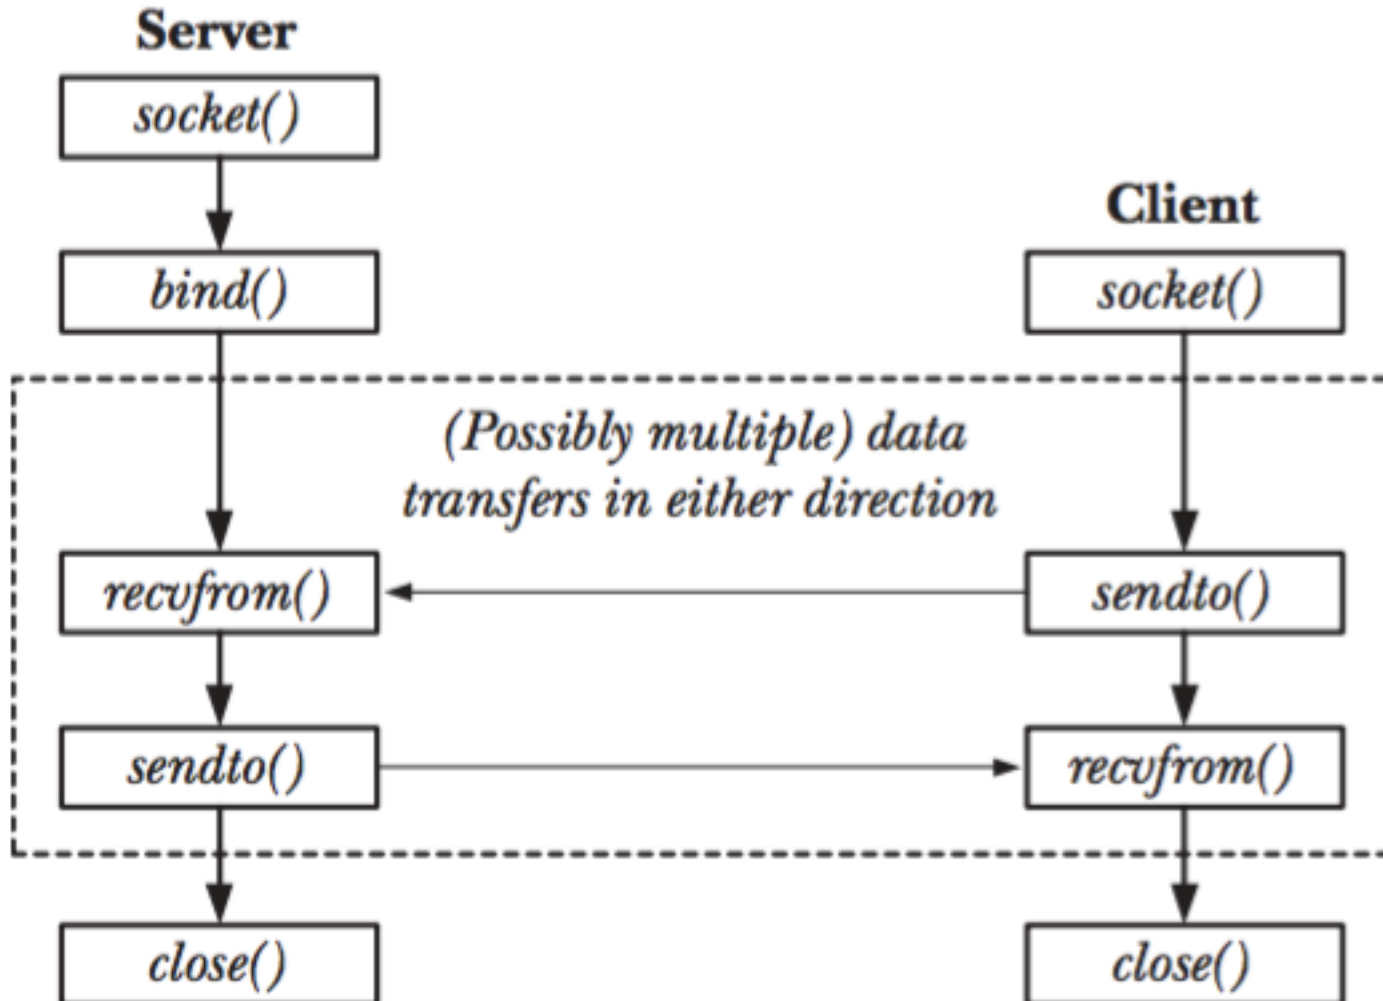


- To perform I/O, we use the *read()* and *write()* system calls
  - Since sockets are bidirectional, both calls may be used on each end of the connection
- A socket may be closed using the *close()* system call or as a consequence of the application terminating

# Datagram Sockets

- Explained by analogy with the postal system:
  - ① *socket()* is setting up a mailbox
  - ② *bind()* is associating the mailbox with my address
    - a client sends datagrams (letters) to that address
  - ③ *sendto()* is putting the recipient's address on a letter and posting it
  - ④ *recvfrom()* is receiving a datagram (letter)

# Overview of System Calls used with Datagram Sockets





# Exchanging Datagrams: `recvfrom()`, `sendto()`

```
#include <sys/socket.h>
```

```
ssize_t recvfrom(int sockfd, void *buffer, size_t length, int flags,  
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

Returns number of bytes received, 0 on EOF, or -1 on error

```
ssize_t sendto(int sockfd, const void *buffer, size_t length, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

Returns number of bytes sent, or -1 on error

- The first three arguments are the same as for *read()* and *write()*
- *flags*: a bit mask controlling socket-specific I/O features (will come to this)
- *src\_addrs*, *dest\_addr*: peer address

# UNIX Domain Sockets

- communication between processes on the same host system
- UNIX Domain Socket Addresses:

```
struct sockaddr_un {  
    sa_family_t sun_family;  
    char sun_path[108];  
};
```

# Binding a UNIX Domain Socket

- initialize a *sockaddr\_un* structure, and then pass a (cast) pointer to this structure as the *addr* argument to *bind()*

---

```
const char *SOCKNAME = "/tmp/mysock";
int sfd;
struct sockaddr_un addr;

sfd = socket(AF_UNIX, SOCK_STREAM, 0);           /* Create socket */
if (sfd == -1)
    errExit("socket");

memset(&addr, 0, sizeof(struct sockaddr_un));    /* Clear structure */
addr.sun_family = AF_UNIX;                       /* UNIX domain address */
strncpy(addr.sun_path, SOCKNAME, sizeof(addr.sun_path) - 1);

if (bind(sfd, (struct sockaddr *) &addr, sizeof(struct sockaddr_un)) == -1)
    errExit("bind");
```

---

# Binding a UNIX Domain Socket

- `memset()`
  - ensures that all of the structure fields have the value 0
- `strncpy()`
  - copies “**SOCKNAME**” to *sun\_path*
  - specifies one less than the size of the *sun\_path* field ensuring that this field always has a terminating null byte

# Binding a UNIX Domain Socket

- Usual to bind a socket to an *absolute* pathname
- A pathname can be bound to *only one* socket
- can't use *open()* to open a socket
- When the socket is no longer required, its pathname entry can be removed using *unlink()* or *remove()*

# Stream Sockets in the UNIX Domain

- A simple client-server application that uses stream sockets in the UNIX domain
  - client: connects to the server, and uses the connection to transfer data from its standard input to the server
  - server: accepts client connections, and transfers all data sent on the connection by the client to standard output

---

```
#include <sys/un.h>
#include <sys/socket.h>
#include "tlpi_hdr.h"

#define SV_SOCKET_PATH "/tmp/us_xfr"

#define BUF_SIZE 100
```

---

```
#include "us_xfr.h"
```

sockets/us\_xfr\_sv.c

```
#define BACKLOG 5
```

```
int
```

```
main(int argc, char *argv[])
```

```
{
```

```
    struct sockaddr_un addr;
```

```
    int sfd, cfd;
```

```
    ssize_t numRead;
```

```
    char buf[BUF_SIZE];
```

```
    sfd = socket(AF_UNIX, SOCK_STREAM, 0);
```

```
    if (sfd == -1)
```

```
        errExit("socket");
```

```
    /* Construct server socket address, bind socket to it,  
       and make this a listening socket */
```

```
    if (remove(SV_SOCKET_PATH) == -1 && errno != ENOENT)
```

```
        errExit("remove-%s", SV_SOCKET_PATH);
```

/\* Remove any existing file with  
the same pathname as that to  
which we bind the socket \*/

```
    memset(&addr, 0, sizeof(struct sockaddr_un));
```

```
    addr.sun_family = AF_UNIX;
```

```
    strncpy(addr.sun_path, SV_SOCKET_PATH, sizeof(addr.sun_path) - 1);
```

```
    if (bind(sfd, (struct sockaddr *) &addr, sizeof(struct sockaddr_un)) == -1)
```

```
        errExit("bind");
```

```
    if (listen(sfd, BACKLOG) == -1)
```

```
        errExit("listen");
```

```
for (;;) {          /* Handle client connections iteratively */

    /* Accept a connection. The connection is returned on a new
       socket, 'cfd'; the listening socket ('sfd') remains open
       and can be used to accept further connections. */

    cfd = accept(sfd, NULL, NULL);
    if (cfd == -1)
        errExit("accept");

    /* Transfer data from connected socket to stdout until EOF */

    while ((numRead = read(cfd, buf, BUF_SIZE)) > 0)
        if (write(STDOUT_FILENO, buf, numRead) != numRead)
            fatal("partial/failed write");

    if (numRead == -1)
        errExit("read");

    if (close(cfd) == -1)
        errMsg("close");

}
}
```



```
#include "us_xfr.h"
```

sockets/us\_xfr\_cl.c

```
int
main(int argc, char *argv[])
{
    struct sockaddr_un addr;
    int sfd;
    ssize_t numRead;
    char buf[BUF_SIZE];

    sfd = socket(AF_UNIX, SOCK_STREAM, 0);      /* Create client socket */
    if (sfd == -1)
        errExit("socket");

    /* Construct server address, and make the connection */

    memset(&addr, 0, sizeof(struct sockaddr_un));
    addr.sun_family = AF_UNIX;
    strncpy(addr.sun_path, SV_SOCK_PATH, sizeof(addr.sun_path) - 1);

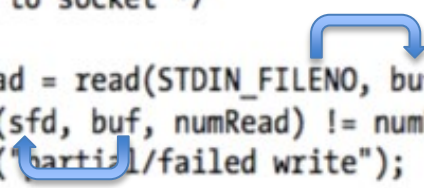
    if (connect(sfd, (struct sockaddr *) &addr,
        sizeof(struct sockaddr_un)) == -1)
        errExit("connect");

    /* Copy stdin to socket */

    while ((numRead = read(STDIN_FILENO, buf, BUF_SIZE)) > 0)
        if (write(sfd, buf, numRead) != numRead)
            fatal("partial/failed write");

    if (numRead == -1)
        errExit("read");

    exit(EXIT_SUCCESS);      /* Closes our socket; server sees EOF */
}
```



# Example Code

```
$ ./us_xfr_sv > b &  
[1] 9866  
$ ls -lF /tmp/us_xfr  
srwxr-xr-x  1 mtk      users      0 Jul 18 10:48 /tmp/us_xfr=
```

*Examine socket file with ls*

We then create a test file to be used as input for the client, and run the client:

```
$ cat *.c > a  
$ ./us_xfr_cl < a
```

*Client takes input from test file*

At this point, the child has completed. Now we terminate the server as well, and check that the server's output matches the client's input:

```
$ kill %1  
[1]+  Terminated    ./us_xfr_sv >b  
$ diff a b  
$
```

*Terminate server*

*Shell sees server's termination*

# Datagram Sockets in the UNIX Domain

- On Linux, you can send quite large datagram
  - Controlled via the `SO_SNDBUF` and various `/proc` files
  - Some other UNIX impose lower limits, 2048 bytes

---

```
sockets/ud_ucase.h
#include <sys/un.h>
#include <sys/socket.h>
#include <ctype.h>
#include "tlpi_hdr.h"

#define BUF_SIZE 10          /* Maximum size of messages exchanged
                             between client to server */

#define SV_SOCK_PATH "/tmp/ud_ucase"
```

```
#include "ud_ucose.h"
```

sockets/ud\_ucose\_sv.c

```
int
main(int argc, char *argv[])
{
    struct sockaddr_un svaddr, claddr;
    int sfd, j;
    ssize_t numBytes;
    socklen_t len;
    char buf[BUF_SIZE];

    sfd = socket(AF_UNIX, SOCK_DGRAM, 0);           /* Create server socket */
    if (sfd == -1)
        errExit("socket");

    /* Construct well-known address and bind server socket to it */

    if (remove(SV_SOCK_PATH) == -1 && errno != ENOENT)
        errExit("remove-%s", SV_SOCK_PATH);

    memset(&svaddr, 0, sizeof(struct sockaddr_un));
    svaddr.sun_family = AF_UNIX;
    strncpy(svaddr.sun_path, SV_SOCK_PATH, sizeof(svaddr.sun_path) - 1);

    if (bind(sfd, (struct sockaddr *) &svaddr, sizeof(struct sockaddr_un)) == -1)
        errExit("bind");
}
```

```
/* Receive messages, convert to uppercase, and return to client */  
  
for (;;) {  
    len = sizeof(struct sockaddr_un);  
    numBytes = recvfrom(sfd, buf, BUF_SIZE, 0,  
                (struct sockaddr *) &claddr, &len);  
    if (numBytes == -1)  
        errExit("recvfrom");  
  
    printf("Server received %ld bytes from %s\n", (long) numBytes,  
          claddr.sun_path);  
  
    for (j = 0; j < numBytes; j++)  
        buf[j] = toupper((unsigned char) buf[j]);  
  
    if (sendto(sfd, buf, numBytes, 0, (struct sockaddr *) &claddr, len) !=  
        numBytes)  
        fatal("sendto");  
}  
}
```



```
#include "ud_ucose.h"
```

sockets/ud\_ucose\_cl.c

```
int
main(int argc, char *argv[])
{
    struct sockaddr_un svaddr, claddr;
    int sfd, j;
    size_t msgLen;
    ssize_t numBytes;
    char resp[BUF_SIZE];

    if (argc < 2 || strcmp(argv[1], "--help") == 0)
        usageErr("%s msg...\n", argv[0]);

    /* Create client socket; bind to unique pathname (based on PID) */

    sfd = socket(AF_UNIX, SOCK_DGRAM, 0);
    if (sfd == -1)
        errExit("socket");

    memset(&claddr, 0, sizeof(struct sockaddr_un));
    claddr.sun_family = AF_UNIX;
    snprintf(claddr.sun_path, sizeof(claddr.sun_path),
             "/tmp/ud_ucose_cl.%ld", (long) getpid());

    if (bind(sfd, (struct sockaddr *) &claddr, sizeof(struct sockaddr_un)) == -1)
        errExit("bind");
```

```

/* Construct address of server */

memset(&svaddr, 0, sizeof(struct sockaddr_un));
svaddr.sun_family = AF_UNIX;
strncpy(svaddr.sun_path, SV_SOCKET_PATH, sizeof(svaddr.sun_path) - 1);

/* Send messages to server; echo responses on stdout */

for (j = 1; j < argc; j++) {
    msgLen = strlen(argv[j]);          /* May be longer than BUF_SIZE */
    if (sendto(sfd, argv[j], msgLen, 0, (struct sockaddr *) &svaddr,
               sizeof(struct sockaddr_un)) != msgLen)
        fatal("sendto");

    numBytes = recvfrom(sfd, resp, BUF_SIZE, 0, NULL, NULL);
    if (numBytes == -1)
        errExit("recvfrom");
    printf("Response %d: %.*s\n", j, (int) numBytes, resp);
}

remove(claddr.sun_path);               /* Remove client socket pathname */
exit(EXIT_SUCCESS);
}

```

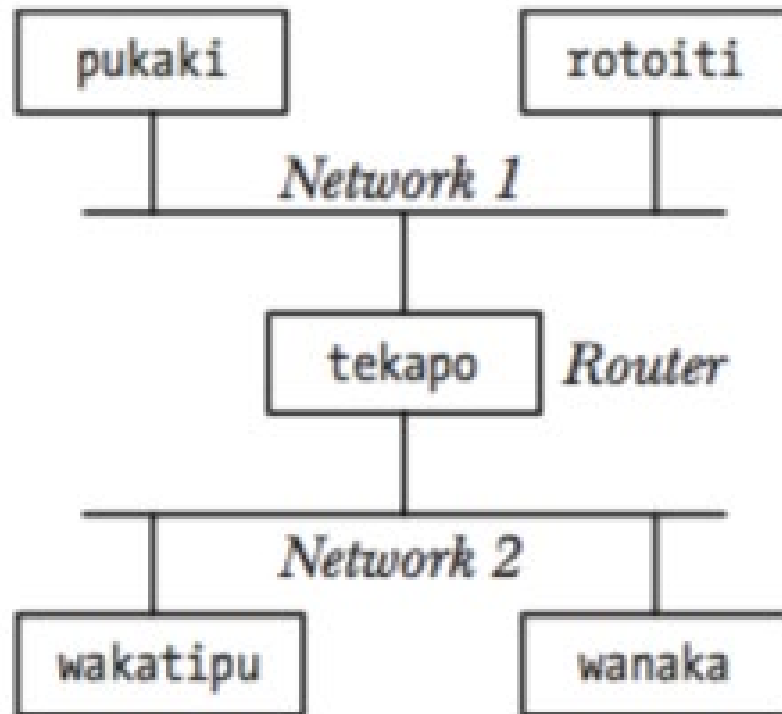
# Results

```
$ ./ud_ucose_sv &  
[1] 20113  
$ ./ud_ucose_cl hello world Send 2 messages to server  
Server received 5 bytes from /tmp/ud_ucose_cl.20150  
Response 1: HELLO  
Server received 5 bytes from /tmp/ud_ucose_cl.20150  
Response 2: WORLD  
$ ./ud_ucose_cl 'long message' Send 1 longer message to server  
Server received 10 bytes from /tmp/ud_ucose_cl.20151  
Response 1: LONG MESSA  
$ kill %1 Terminate server
```

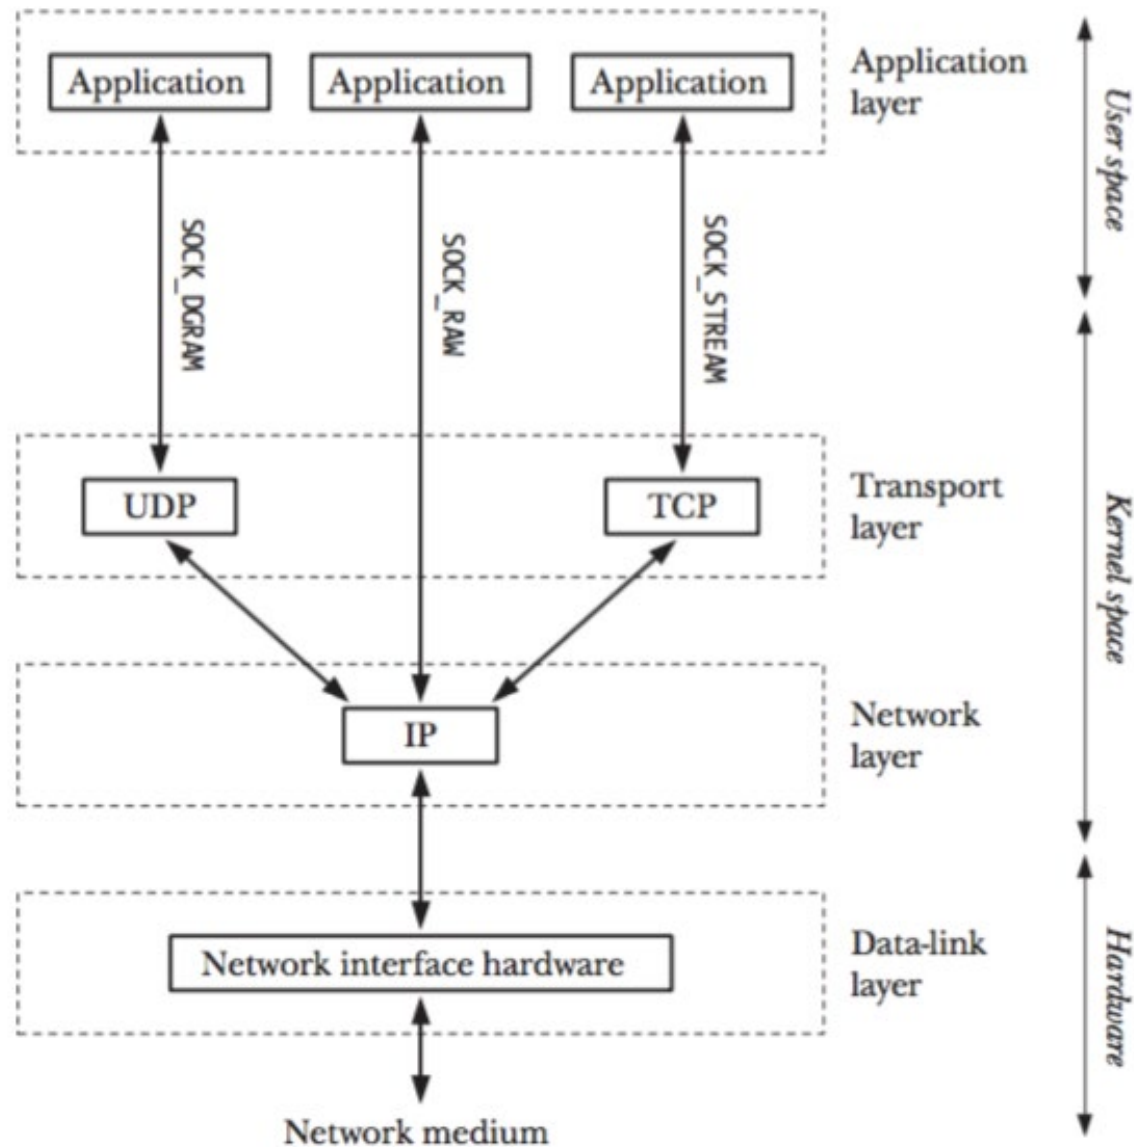


# Sockets: TCP/IP Networks

- TCP/IP is dominant protocol



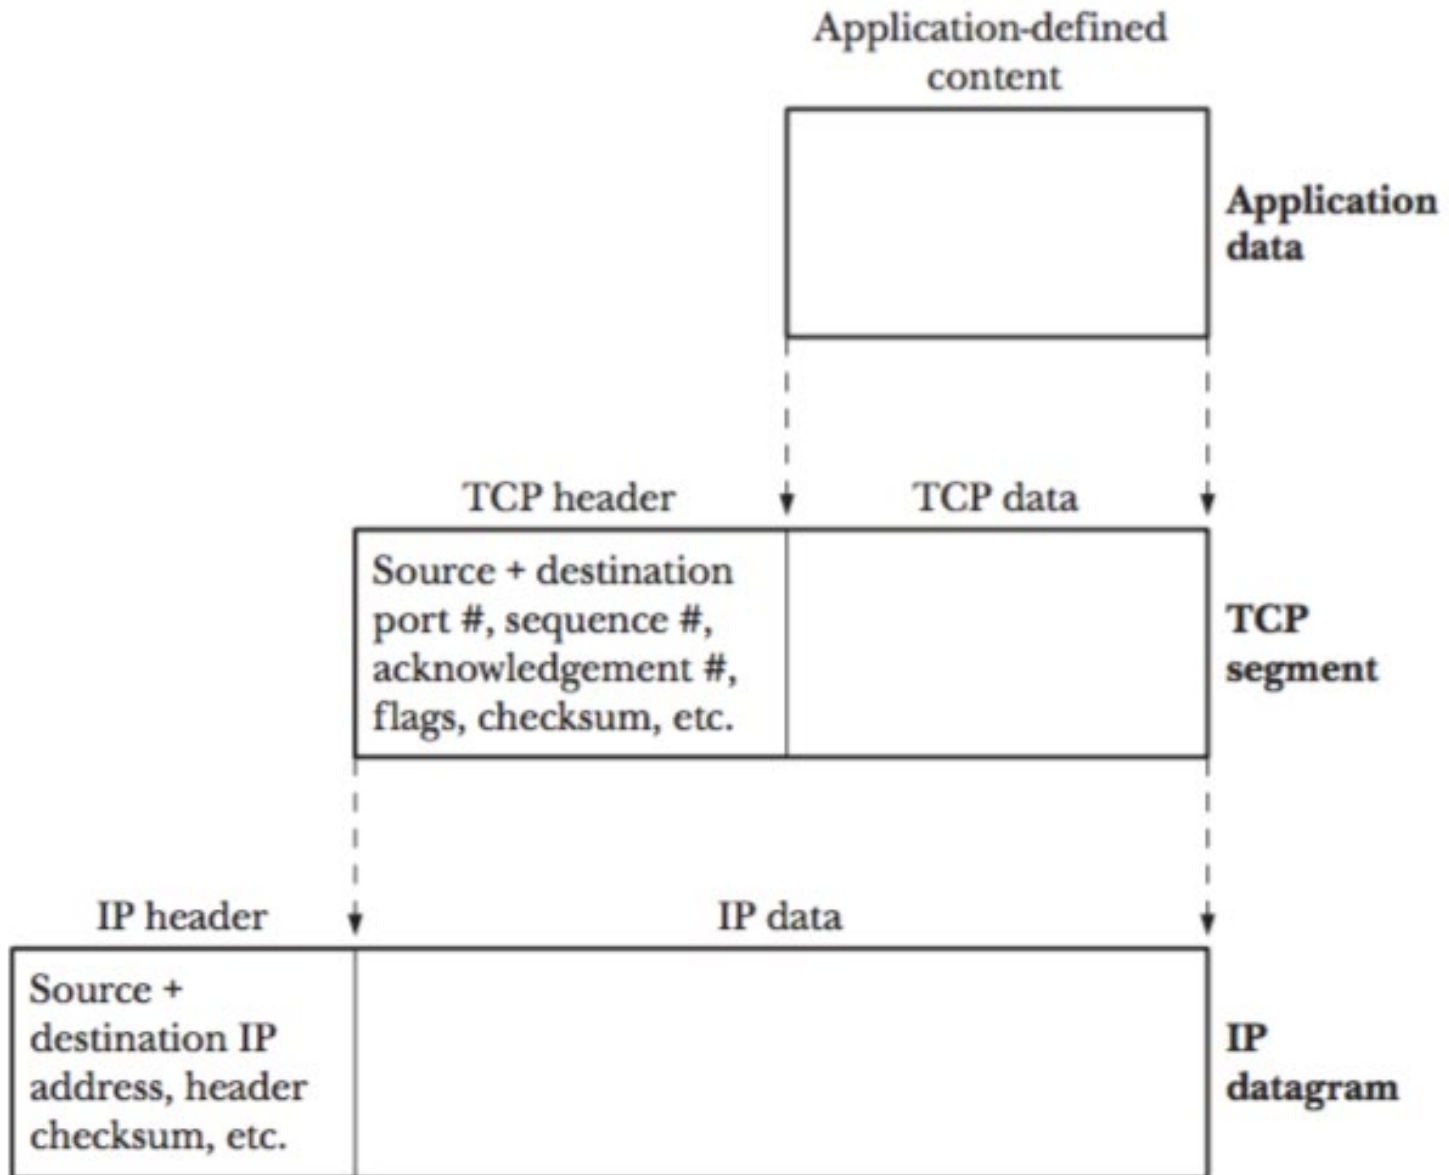
# Networking Protocol and Layers



# Encapsulation

- The lower layer makes no attempt to interpret information sent from the upper layer
  - and adds its own layer-specific header before passing the packet down to the next lower layer
- When data is passed up from a lower layer to a higher layer, a converse unpacking process takes place

# Encapsulation




# Data-Link Layer

- consists of device driver and H/W interface (network card)
- concerned with transferring data across a physical link in a network
- encapsulates datagrams into frames
- includes a header containing destination address and frame size
- performs error detection, retransmission, flow control
  - Some data-link layers split large network packets into multiple frames and reassemble them at the receiver

# Network Layer: IP

- concerned with delivering packets from source to destination host
- Tasks
  - breaks data into fragments small enough for transmission
  - routes data across the internet
  - provides services to the transport layer
- IP transmits data in the form of datagram (packets)
- Header contains
  - address of the target host
  - originating address of the packet

# IP Address



- Network ID
    - specifies the network on which a host resides
  - Host ID
    - identifies the host within that network
  - network mask
    - 1s indicates which part of the address contains the assigned network ID
    - 0s to assign unique host IDs on its network
- e.g., 204.152.189.0/24 : /24 indicates that network ID part of the assigned address consists of the leftmost 24 bits (network mask is 255.255.255.0)

# IP Address

- Loopback address
  - special address 127.0.0.1 : conventionally assigned to the hostname localhost
  - A datagram sent to this address never actually reaches the network for testing client/server programs on the same host
  - **INADDR\_LOOPBACK** is defined for this address
- Wildcard address
  - INADDR\_ANY
  - useful for applications that bind Internet domain socket on multi-homed hosts (more than one network interface)



# Transport Layer

- Two widely used protocols
  - UDP (User Datagram Protocol) : for datagram sockets
  - TCP (Transmission Control Protocol) : for stream sockets
- Task is to provide an end-to-end communication service to applications residing on different hosts
  - requires a method to differentiate the applications
  - this differentiation is provided by a 16-bit *port* number

# Port Numbers

- Well-known port numbers
  - ssh (secure shell) : 22
  - HTTP : 80
  - HTTPS: 443
  - assigned 0 ~ 1023
  - Normally privileged (**CAP\_NET\_BIND\_SERVICE**)
- IANA (Internet Assigned Numbers Authority) records registered ports
  - 1024 ~ 49151

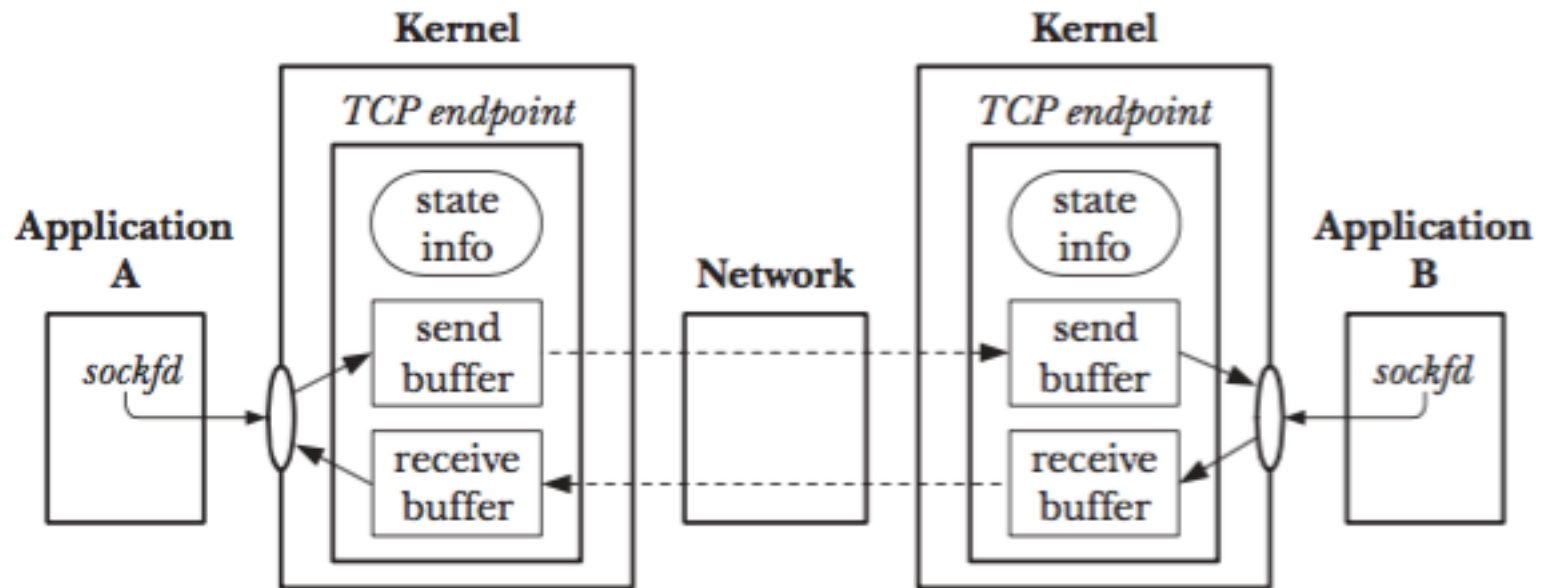
# UDP

- adds two features
  - port number
  - data checksum (since connectionless)
- data checksum
  - 16 bits long
  - for error detection

# TCP

- Acknowledgements
  - When a TCP segment arrives at its destination without errors, the receiving TCP sends a positive acknowledgement to the sender
- Sequencing
  - Each byte transmitted over a TCP is assigned a logical sequence number
  - This number indicates the position in data stream

# TCP



# Network Byte Order

- IP addresses and port numbers are integer values
- When passing these numbers across a network, different architectures store integer bytes differently
- Thus, standard ordering must be used
  - this ordering is called network byte order
  - can be stored in a *socket address structure*

# Network Byte Order

- Big-endian: stores integers with the most significant byte first
- Little-endian: least significant byte first

2-byte integer		4-byte integer			
		address <i>N</i>	address <i>N + 1</i>	address <i>N + 2</i>	address <i>N + 3</i>
Big-endian byte order		1 (MSB)	0 (LSB)		
				3 (MSB)	2 (LSB)
				1 (MSB)	0 (LSB)
		address <i>N</i>	address <i>N + 1</i>	address <i>N + 2</i>	address <i>N + 3</i>
Little-endian byte order		0 (LSB)	1 (MSB)		
				0 (LSB)	1 (MSB)
				2 (MSB)	3 (LSB)

*MSB = Most Significant Byte, LSB = Least Significant Byte*

# Host Byte Order

- convention of the host machine
- should be converted to network byte order before storing them in socket address structures



# htons(), htonl(), ntohs(), ntohl()

```
#include <arpa/inet.h>
```

```
uint16_t htons(uint16_t host_uint16);
```

Returns *host\_uint16* converted to network byte order

```
uint32_t htonl(uint32_t host_uint32);
```

Returns *host\_uint32* converted to network byte order

```
uint16_t ntohs(uint16_t net_uint16);
```

Returns *net\_uint16* converted to host byte order

```
uint32_t ntohl(uint32_t net_uint32);
```

Returns *net\_uint32* converted to host byte order

- convert between host byte order and network byte order

# readline()

```
#include "read_line.h"
```

```
ssize_t readLine(int fd, void *buffer, size_t n);
```

Returns number of bytes copied into *buffer* (excluding terminating null byte), or 0 on end-of-file, or -1 on error

- reads bytes from the file referred to by the file descriptor, *fd* until a newline is encountered

# Reading data a line at a time

---

sockets/read\_line.c

```
#include <unistd.h>
#include <errno.h>
#include "read_line.h"          /* Declaration of readLine() */

ssize_t
readLine(int fd, void *buffer, size_t n)
{
    ssize_t numRead;            /* # of bytes fetched by last read() */
    size_t totRead;            /* Total bytes read so far */
    char *buf;
    char ch;

    if (n <= 0 || buffer == NULL) {
        errno = EINVAL;
        return -1;
    }

    buf = buffer;               /* No pointer arithmetic on "void *" */

    totRead = 0;
```

```

for (;;) {
    numRead = read(fd, &ch, 1);

    if (numRead == -1) {
        if (errno == EINTR)          /* Interrupted --> restart read() */
            continue;
        else
            return -1;                /* Some other error */

    } else if (numRead == 0) {        /* EOF */
        if (totRead == 0)             /* No bytes read; return 0 */
            return 0;
        else                          /* Some bytes read; add '\0' */
            break;

    } else {                          /* 'numRead' must be 1 if we get here */
        if (totRead < n - 1) {        /* Discard > (n - 1) bytes */
            totRead++;
            *buf++ = ch;
        }

        if (ch == '\n')
            break;
    }
}

*buf = '\0';
return totRead;
}

```

# Internet Socket Addresses

```
struct in_addr {                                /* IPv4 4-byte address */
    in_addr_t s_addr;                          /* Unsigned 32-bit integer */
};

struct sockaddr_in {                            /* IPv4 socket address */
    sa_family_t sin_family;                    /* Address family (AF_INET) */
    in_port_t sin_port;                       /* Port number */
    struct in_addr sin_addr;                  /* IPv4 address */
    unsigned char __pad[X];                  /* Pad to size of 'sockaddr'
                                           structure (16 bytes) */
};
```

- IPv4 socket addresses: *struct sockaddr\_in*
- *sin\_family* field is always set to **AF\_INET**

# Internet Socket Addresses

```
struct in6_addr {                /* IPv6 address structure */
    uint8_t s6_addr[16];        /* 16 bytes == 128 bits */
};


struct sockaddr_in6 {            /* IPv6 socket address */
    sa_family_t sin6_family;     /* Address family (AF_INET6) */
    in_port_t sin6_port;         /* Port number */
    uint32_t sin6_flowinfo;      /* IPv6 flow information */
    struct in6_addr sin6_addr;   /* IPv6 address */
    uint32_t sin6_scope_id;      /* Scope ID (new in kernel 2.4) */
};
```

- IPv6 socket addresses: *struct sockaddr\_in6*
- IPv6 loopback address is (::1)

# inet\_pton(), inet\_ntop()

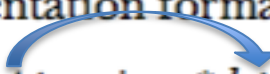
```
#include <arpa/inet.h>
```

```
int inet_pton(int domain, const char *src_str, void *addrp);
```



Returns 1 on successful conversion, 0 if *src\_str* is not in presentation format, or -1 on error

```
const char *inet_ntop(int domain, const void *addrp, char *dst_str, size_t len);
```



Returns pointer to *dst\_str* on success, or NULL on error

- converts between binary form and dotted-decimal notation
- *p* : presentation (human readable)
  - 204.111.122.114 (IPv4 dotted-decimal address)
  - ::1 (an IPv6 colon-separated hexadecimal address)
- *n* : network (binary form)

# Example (Datagram Sockets)

- header file

---

```
sockets/i6d_ucase.h

#include <netinet/in.h>
#include <arpa/inet.h>
#include <sys/socket.h>
#include <ctype.h>
#include "tlpi_hdr.h"

#define BUF_SIZE 10                /* Maximum size of messages exchanged
                                   between client and server */

#define PORT_NUM 50002             /* Server port number */
```



```
#include "i6d_ucose.h"

int
main(int argc, char *argv[])
{
    struct sockaddr_in6 svaddr, claddr;
    int sfd, j;
    ssize_t numBytes;
    socklen_t len;
    char buf[BUF_SIZE];
    char claddrStr[INET6_ADDRSTRLEN];

    sfd = socket(AF_INET6, SOCK_DGRAM, 0);
    if (sfd == -1)
        errExit("socket");

    memset(&svaddr, 0, sizeof(struct sockaddr_in6));
    svaddr.sin6_family = AF_INET6;
    svaddr.sin6_addr = in6addr_any;           /* Wildcard address */
    svaddr.sin6_port = htons(PORT_NUM);

    if (bind(sfd, (struct sockaddr *) &svaddr,
        sizeof(struct sockaddr_in6)) == -1)
        errExit("bind");
```

```

/* Receive messages, convert to uppercase, and return to client */


for (;;) {
    len = sizeof(struct sockaddr_in6);
    numBytes = recvfrom(sfd, buf, BUF_SIZE, 0,
        (struct sockaddr *) &claddr, &len);
    if (numBytes == -1)
        errExit("recvfrom");

    if (inet_ntop(AF_INET6, &claddr.sin6_addr, claddrStr,
        INET6_ADDRSTRLEN) == NULL)
        printf("Couldn't convert client address to string\n");
    else
        printf("Server received %ld bytes from (%s, %u)\n",
            (long) numBytes, claddrStr, ntohs(claddr.sin6_port));

    for (j = 0; j < numBytes; j++)
        buf[j] = toupper((unsigned char) buf[j]);

    if (sendto(sfd, buf, numBytes, 0, (struct sockaddr *) &claddr, len) !=
        numBytes)
        fatal("sendto");
}
}

```



```
#include "i6d_ucase.h"

int
main(int argc, char *argv[])
{
    struct sockaddr_in6 svaddr;
    int sfd, j;
    size_t msglen;
    ssize_t numBytes;
    char resp[BUF_SIZE];

    if (argc < 3 || strcmp(argv[1], "--help") == 0)
        usageErr("%s host-address msg...\n", argv[0]);

    sfd = socket(AF_INET6, SOCK_DGRAM, 0);      /* Create client socket */
    if (sfd == -1)
        errExit("socket");

    memset(&svaddr, 0, sizeof(struct sockaddr_in6));
    svaddr.sin6_family = AF_INET6;
    svaddr.sin6_port = htons(PORT_NUM);
    if (inet_pton(AF_INET6, argv[1], &svaddr.sin6_addr) <= 0)
        fatal("inet_pton failed for address '%s'", argv[1]);
}
```

```

/* Send messages to server; echo responses on stdout */

for (j = 2; j < argc; j++) {
    msgLen = strlen(argv[j]);
    if (sendto(sfd, argv[j], msgLen, 0, (struct sockaddr *) &svaddr,
              sizeof(struct sockaddr_in6)) != msgLen)
        fatal("sendto");

    numBytes = recvfrom(sfd, resp, BUF_SIZE, 0, NULL, NULL);
    if (numBytes == -1)
        errExit("recvfrom");

    printf("Response %d: %.*s\n", j - 1, (int) numBytes, resp);
}

exit(EXIT_SUCCESS);
}

```

```
$ ./i6d_ucose_sv &
```

```
[1] 31047
```

```
$ ./i6d_ucose_cl ::1 ciao
```

*Send to server on local host*

```
Server received 4 bytes from (::1, 32770)
```

```
Response 1: CIAO
```

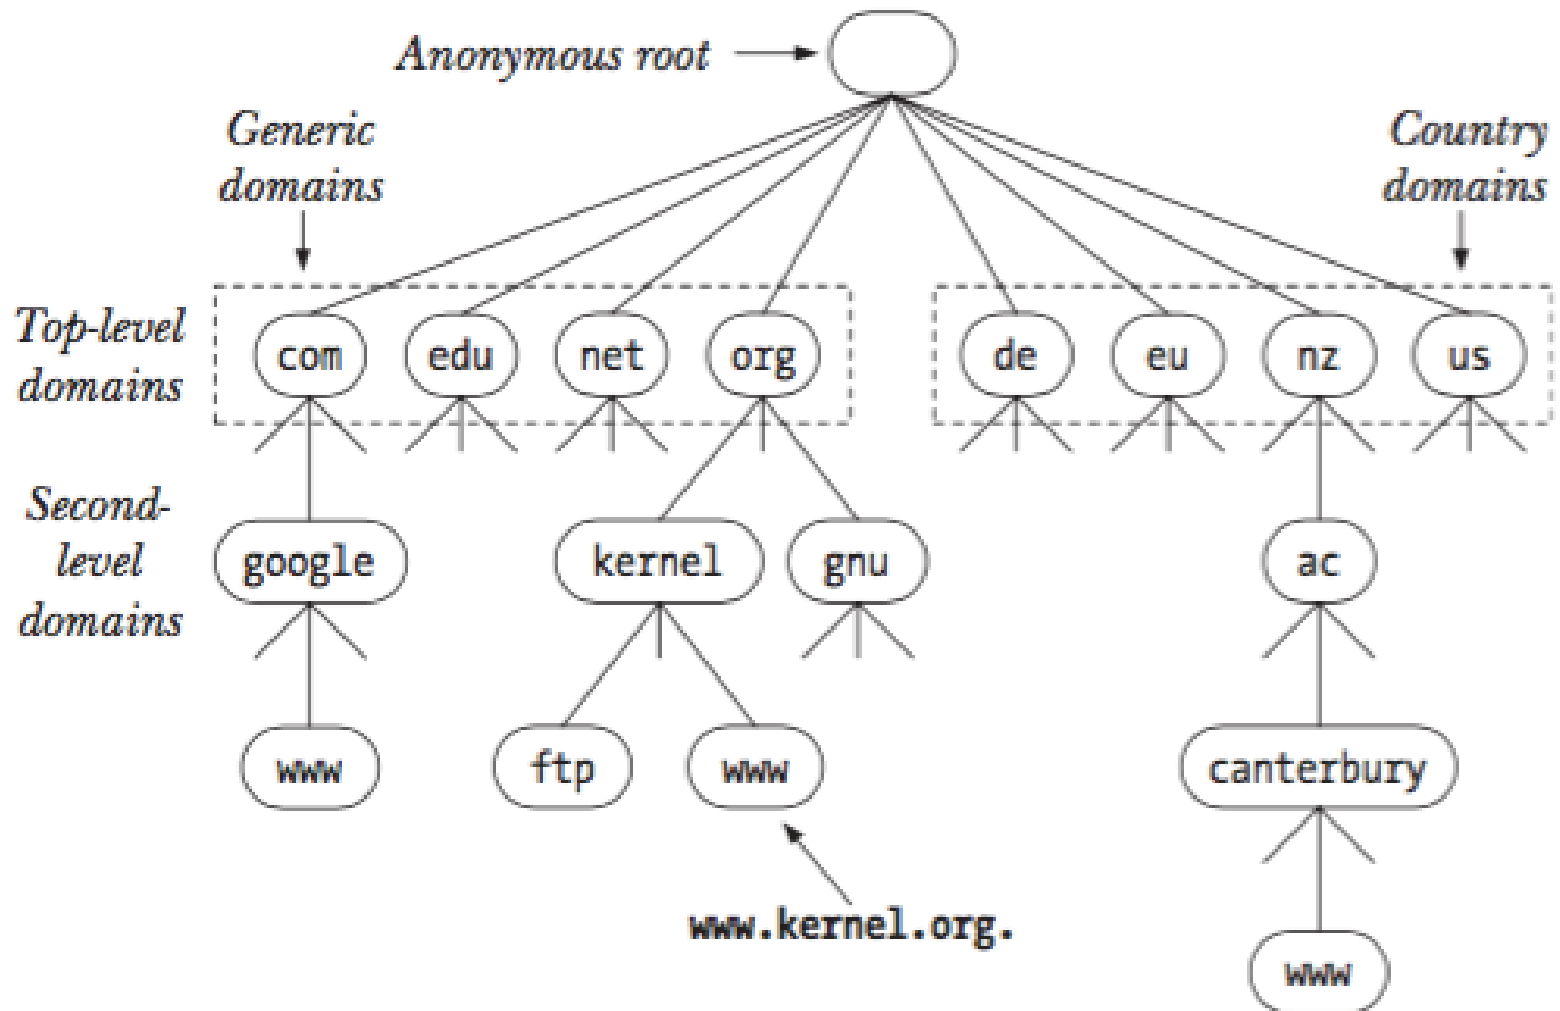
# DNS

```
$ cat /etc/hosts
```

```
# IP-address canonical hostname [aliases]
```

```
127.0.0.1 localhost
```

# A subset of DNS hierarchy



# getaddrinfo()

- converts *host* (e.g., `www.google.com`) and *service* names (e.g., `https`) to *IP* and *port*
  - returns a list of socket address structures, each of which contains an IP address and port number

```
#include <sys/socket.h>
#include <netdb.h>
```

```
int getaddrinfo(const char *host, const char *service,
               const struct addrinfo *hints, struct addrinfo **result);
```

Returns 0 on success, or nonzero on error

- service: either service name or decimal port number
- hints: addrinfo structure
  - that specifies further criteria for selecting the socket address structures
- result: returns a list of structures

# getaddrinfo()

- When a program calls *getaddrinfo()* to obtain the IP address for a domain name, *getaddrinfo()* employs a suite of library functions that communicate with local DNS server
  - if this server can't supply the required information, then it communicates with other DNS servers

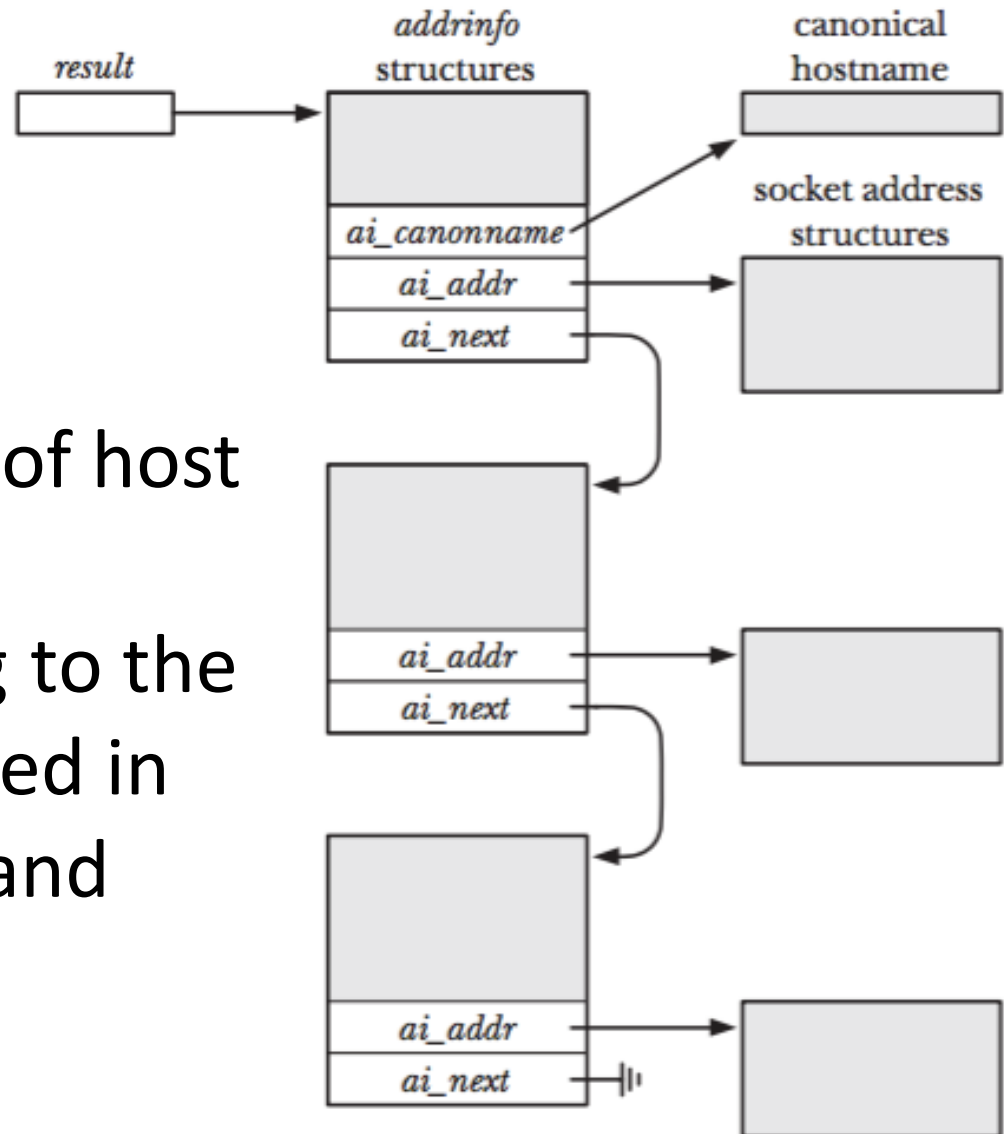


# getaddrinfo()

```
struct addrinfo {  
    int     ai_flags;           /* Input flags (AI_* constants) */  
    int     ai_family;         /* Address family */  
    int     ai_socktype;       /* Type: SOCK_STREAM, SOCK_DGRAM */  
    int     ai_protocol;       /* Socket protocol */  
    size_t  ai_addrlen;        /* Size of structure pointed to by ai_addr */  
    char    *ai_canonname;     /* Canonical name of host */  
    struct sockaddr *ai_addr;  /* Pointer to socket address structure */  
    struct addrinfo *ai_next;  /* Next structure in linked list */  
};
```

# Structures allocated and returned by *getaddrinfo()*

- There may be multiple combinations of host and service corresponding to the criteria specified in *host*, *service*, and *hints*



# getnameinfo()

- converse of getaddrinfo()
- *addr*: a pointer to the socket address structure that is to be converted
- *host*, *service*: resulting host and service names returned in the buffers pointed to by

```
#include <sys/socket.h>
#include <netdb.h>
```

```
int getnameinfo(const struct sockaddr *addr, socklen_t addrlen, char *host,
                size_t hostlen, char *service, size_t servlen, int flags);
```

Returns 0 on success, or nonzero on error

# Example (Stream Sockets)

---

```
sockets/is_seqnum.h
#include <netinet/in.h>
#include <sys/socket.h>
#include <signal.h>
#include "read_line.h"      /* Declaration of readLine() */
#include "tlpi_hdr.h"

#define PORT_NUM "50000"    /* Port number for server */

#define INT_LEN 30          /* Size of string able to hold largest
                             integer (including terminating '\n') */
```

---

sockets/is\_seqnum.h

```

#define _BSD_SOURCE          /* To get definitions of NI_MAXHOST and
                             NI_MAXSERV from <netdb.h> */

#include <netdb.h>
#include "is_seqnum.h"

#define BACKLOG 50

int
main(int argc, char *argv[])
{
    uint32_t seqNum;
    char reqLenStr[INT_LEN];          /* Length of requested sequence */
    char seqNumStr[INT_LEN];          /* Start of granted sequence */
    struct sockaddr_storage claddr;
    int lfd, cfd, optval, reqLen;
    socklen_t addrlen;
    struct addrinfo hints;
    struct addrinfo *result, *rp;
#define ADDRSTRLEN (NI_MAXHOST + NI_MAXSERV + 10)
    char addrStr[ADDRSTRLEN];
    char host[NI_MAXHOST];
    char service[NI_MAXSERV];

    if (argc > 1 && strcmp(argv[1], "--help") == 0)
        usageErr("%s [init-seq-num]\n", argv[0]);

```

① seqNum = (argc > 1) ? getInt(argv[1], 0, "init-seq-num") : 0;

```

② if (signal(SIGPIPE, SIG_IGN) == SIG_ERR)
    errExit("signal");

/* Call getaddrinfo() to obtain a list of addresses that
   we can try binding to */

memset(&hints, 0, sizeof(struct addrinfo));
hints.ai_canonname = NULL;
hints.ai_addr = NULL;
hints.ai_next = NULL;
hints.ai_socktype = SOCK_STREAM;
hints.ai_family = AF_UNSPEC;          /* Allows IPv4 or IPv6 */
③ hints.ai_flags = AI_PASSIVE | AI_NUMERICSERV;
                        /* Wildcard IP address; service name is numeric */
④ if (getaddrinfo(NULL, PORT_NUM, &hints, &result) != 0)
    errExit("getaddrinfo");

/* Walk through returned list until we find an address structure
   that can be used to successfully create and bind a socket */

optval = 1;
⑤ for (rp = result; rp != NULL; rp = rp->ai_next) {
    lfd = socket(rp->ai_family, rp->ai_socktype, rp->ai_protocol);
    if (lfd == -1)
        continue;                      /* On error, try next address */

```

```

⑥      if (setsockopt(lfd, SOL_SOCKET, SO_REUSEADDR, &optval, sizeof(optval))
          == -1)
          errExit("setsockopt");

⑦      if (bind(lfd, rp->ai_addr, rp->ai_addrlen) == 0)
          break; /* Success */

          /* bind() failed: close this socket and try next address */

          close(lfd);
      }

      if (rp == NULL)
          fatal("Could not bind socket to any address");

⑧      if (listen(lfd, BACKLOG) == -1)
          errExit("listen");

          freeaddrinfo(result);

⑨      for (;;) {                                /* Handle clients iteratively */

          /* Accept a client connection, obtaining client's address */

          addrlen = sizeof(struct sockaddr_storage);
          cfd = accept(lfd, (struct sockaddr *) &claddr, &addrlen);
          if (cfd == -1) {
              errMsg("accept");
              continue;
          }

```

```

⑪ if (getnameinfo((struct sockaddr *)&claddr, addrlen,
    host, NI_MAXHOST, service, NI_MAXSERV, 0) == 0)
    snprintf(addrStr, ADDRSTRLEN, "(%s, %s)", host, service);
else
    snprintf(addrStr, ADDRSTRLEN, "(?UNKNOWN?)");
printf("Connection from %s\n", addrStr);

/* Read client request, send sequence number back */

⑫ if (readLine(cfd, reqLenStr, INT_LEN) <= 0) {
    close(cfd);
    continue; /* Failed read; skip request */
}

⑬ reqLen = atoi(reqLenStr);
if (reqLen <= 0) { /* Watch for misbehaving clients */
    close(cfd);
    continue; /* Bad request; skip it */
}

⑭ snprintf(seqNumStr, INT_LEN, "%d\n", seqNum);
if (write(cfd, &seqNumStr, strlen(seqNumStr)) != strlen(seqNumStr))
    fprintf(stderr, "Error on write");

⑮ seqNum += reqLen; /* Update sequence number */

if (close(cfd) == -1) /* Close connection */
    errMsg("close");
}
}

```



```
#include <netdb.h>
#include "is_seqnum.h"

int
main(int argc, char *argv[])
{
    char *reqLenStr;                /* Requested length of sequence */
    char seqNumStr[INT_LEN];        /* Start of granted sequence */
    int cfd;
    ssize_t numRead;
    struct addrinfo hints;
    struct addrinfo *result, *rp;

    if (argc < 2 || strcmp(argv[1], "--help") == 0)
        usageErr("%s server-host [sequence-len]\n", argv[0]);

    /* Call getaddrinfo() to obtain a list of addresses that
       we can try connecting to */

    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_canonname = NULL;
    hints.ai_addr = NULL;
    hints.ai_next = NULL;
    hints.ai_family = AF_UNSPEC;    /* Allows IPv4 or IPv6 */
    hints.ai_socktype = SOCK_STREAM;
    hints.ai_flags = AI_NUMERICSERV;

    ① if (getaddrinfo(argv[1], PORT_NUM, &hints, &result) != 0)
        errExit("getaddrinfo");
```

```
/* Walk through returned list until we find an address structure  
that can be used to successfully connect a socket */
```

```
② for (rp = result; rp != NULL; rp = rp->ai_next) {  
③   cfd = socket(rp->ai_family, rp->ai_socktype, rp->ai_protocol);  
   if (cfd == -1)  
       continue;                               /* On error, try next address */  
  
④   if (connect(cfd, rp->ai_addr, rp->ai_addrlen) != -1)  
       break;                                   /* Success */  
  
   /* Connect failed: close this socket and try next address */  
  
   close(cfd);  
}  
  
if (rp == NULL)  
    fatal("Could not connect socket to any address");  
  
freeaddrinfo(result);  
  
/* Send requested sequence length, with terminating newline */  
  
⑤ reqLenStr = (argc > 2) ? argv[2] : "1";  
if (write(cfd, reqLenStr, strlen(reqLenStr)) != strlen(reqLenStr))  
    fatal("Partial/failed write (reqLenStr)");  
if (write(cfd, "\n", 1) != 1)  
    fatal("Partial/failed write (newline)");
```

```
/* Read and display sequence number returned by server */
```

```
⑥ numRead = readLine(cfd, seqNumStr, INT_LEN);  
  if (numRead == -1)  
      errExit("readLine");  
  if (numRead == 0)  
      fatal("Unexpected EOF from server");  
  
⑦ printf("Sequence number: %s", seqNumStr);          /* Includes '\n' */  
  
  exit(EXIT_SUCCESS);                                /* Closes 'cfd' */  
}  
  
_____ sockets/is_seqnum_cl.c
```

# Results

```
$ ./is_seqnum_sv &
```

```
[1] 4075
```

```
$ ./is_seqnum_cl localhost
```

```
Connection from (localhost, 33273)
```

```
Sequence number: 0
```

```
$ ./is_seqnum_cl localhost 10
```

```
Connection from (localhost, 33274)
```

```
Sequence number: 1
```

```
$ ./is_seqnum_cl localhost
```

```
Connection from (localhost, 33275)
```

```
Sequence number: 11
```

*Client 1: requests 1 sequence number*

*Server displays client address + port*

*Client displays returned sequence number*

*Client 2: requests 10 sequence numbers*

*Client 3: requests 1 sequence number*

```
$ telnet localhost 50000
```

*Our server uses this port number*

*Empty line printed by telnet*

```
Trying 127.0.0.1...
```

```
Connection from (localhost, 33276)
```

```
Connected to localhost.
```

```
Escape character is '^]'.
```

```
1
```

```
12
```

```
Connection closed by foreign host.
```

*Enter length of requested sequence*

*telnet displays sequence number and*

*detects that server closed connection*